

Uber vs. Taxi: A Driver's Eye View*

Joshua D. Angrist	Sydnee Caldwell	Jonathan V. Hall
MIT, IZA and NBER	Microsoft Research	Uber Technologies, Inc.

July 15, 2020

Abstract

Rideshare drivers pay a proportion of their fares to a ride-hailing platform operator, a commission-based compensation model used by many service providers. To Uber drivers, this commission is known as the Uber fee. By contrast, traditional taxi drivers in most US cities make a fixed payment independent of their earnings, usually a weekly or daily medallion lease, keeping every fare dollar net of lease costs and other expenses. We assess these compensation models using an experiment that offered random samples of Boston Uber drivers opportunities to lease a virtual taxi medallion that eliminates the Uber fee. Some drivers were offered a negative fee. Drivers' labor supply response to our offers reveals a large intertemporal substitution elasticity, on the order of 1.2, and higher for those who accept lease contracts. At the same time, our virtual lease program was under-subscribed: many drivers who would have benefitted from buying an inexpensive lease chose to sit out. We use these results to compute the average compensation required to make drivers indifferent between rideshare and taxi-style compensation contracts. The results suggest that rideshare drivers gain considerably from the opportunity to drive without leasing.

*Angrist: MIT, IZA and NBER, angrist@mit.edu. Caldwell: Microsoft Research, scaldwell@berkeley.edu. Hall: Uber Technologies, Inc. jvh@uber.com. Special thanks go to Emily Oehlsen for indispensable project management, to the staff of Uber's Boston city team for hosting and assisting with our study, and to Phoebe Cai, Anran Li, Gina Li and Yulu Tang for research assistance. Thanks also go to seminar participants at the Bank of Italy, Chicago, Harvard, The IDC (Israel), MIT, NBER Labor Studies, Nova (Lisbon), Princeton, the Stanford Graduate School of Business, Tinbergen, Uber, U Tokyo, and Uppsala and to Daron Acemoglu, David Card, Liran Einav, Amy Finkelstein, Henry Farber, Ed Glaeser, Nathan Hendren, Guido Imbens, Michael Ostrovsky, Amanda Pallais, and Frank Schilbach for helpful discussions and comments. The views expressed here are those of the authors and do not necessarily reflect those of Uber Technologies, Inc. or MIT. Angrist and Caldwell's work on this project was carried out under a data use agreement executed between MIT and Uber. This study is registered in the AEA RCT Registry as trial no. AEARCTR-0001656. Caldwell thanks the National Science Foundation for support under NSF Doctoral Dissertation Research in Economics award No. 1729822. This is a revised version of NBER Working Paper 23891. The project was approved by MIT's Committee on the Use of Human Subjects (COUHES) protocol #1605555831.

Other Driver issues [the New York Taxi and Limousine Commission] identified include the perceived inflexibility of leases currently offered by lessors as well as the stress associated with starting shifts “in the red” having paid a set lease price at the beginning of shifts—2015 New York Taxi and Limousine Commission Resolution

1 Introduction

Traditional taxi drivers in most large American cities must own or lease one of a limited number of medallions granting them the right to drive. Until recently, limited supply had turned medallions into valuable assets, typically held by investors or fleet owners, and trading for hundreds of thousands of dollars. Most big city taxi drivers therefore lease their medallions by the shift, day, or week. Taxi drivers can drive as much or as little as they want, but they’re on the hook for the lease. The rise of rideshare platforms, including Uber, means that many workers now have the opportunity to add to their earnings by driving private vehicles, no medallion lease required. By the summer of 2016, Uber had almost 20,000 active drivers in Boston, a figure that can be compared with Boston’s long-fixed 1,825 taxi medallions.

In addition to reducing entry barriers and perhaps taxi fares, an important feature of the rideshare model is a proportional compensation scheme, with few fixed costs.¹ In return for a percentage of their earnings known to drivers as a fee or commission, rideshare drivers can set a work schedule without having to worry about covering a lease. Drivers who work long hours are still better off leasing because they keep every dollar earned on a relatively high farebox. But drivers with low hours should prefer to work on a rideshare platform.

This paper uses a series of randomized experiments conducted at a national rideshare company to compare the value of the proportional compensation scheme offered by rideshare companies with traditional taxi compensation (Angrist et al., 2020). The latter can be seen as an exemplar of work arrangements whereby workers “buy the firm” in the sense that they keep every dollar earned after expenses. Our experiments offered random samples of Boston Uber drivers the opportunity to buy a virtual lease that eliminated or reduced the Uber fee. Some lease-paying drivers were offered a negative fee, capturing a possibly higher taxi wage.

¹Some cities, including New York and (until recently) Houston, impose additional licensure and training costs on ride-hailing drivers.

We use driver’s labor supply behavior and contract choices to estimate the parameters that determine the value of a rideshare compensation contract. The first key parameter of interest is the labor supply response to temporarily higher wage rates, or intertemporal substitution elasticity (ISE). A large ISE tends to make medallion-type contracts more attractive because elastic drivers collect additional surplus by driving longer hours when their hourly wage goes up. Drivers’ response to experimental wage changes reveals an ISE for the wage effect on Uber hours, on the order of 1.2 overall, and around 1.8 for drivers who lease. These estimates are broadly consistent with experimental estimates reported for Swiss bicycle messengers by Fehr and Goette (2007).² Our estimated substitution elasticities are remarkably stable across groups with varying levels of experience and work intensity. They’re also broadly in line with the Mas and Pallais (2018) experimental estimates of compensated elasticities for part-time workers who work flexible hours.

The second key parameter in our framework quantifies the extent to which attractive leasing arrangements were under-subscribed. Many drivers to whom we offered a lease indeed took it. But many drivers who would have benefitted from leasing failed to take advantage of the opportunity to do so. We refer to this behavior as “lease aversion,” and use a model of context-specific loss aversion to explain it. Specifically, we compute a behavioral lease-aversion coefficient that rationalizes the lease take-up rates seen in our experiment. This coefficient is also reasonably stable across driver groups. Even without lease aversion, a switch from leasing to proportional compensation generates considerable surplus for most of the drivers in our sample unless lease prices fall below about \$100 per week. In the face of a \$200 weekly lease, lease aversion increases Boston drivers’ average rideshare surplus to nearly one-third of their Uber earnings.

The proportional compensation scheme is not the only difference between rideshare and traditional taxi. Prior work has shown that rideshare drivers—especially female and low-income drivers—value the ability to drive flexible shifts, with no minimum shift requirement (Chen et al., Forthcoming, 2020). Some drivers may benefit from wages that quickly respond to supply and demand (“surge” pricing) (Castillo, 2019) or may value driving their own vehicle. However, our results may explain why workers rarely “buy the firm”.

Price theory suggests that lease-type arrangements are more efficient than a proportional

²See Farber (2005; 2015) for more on taxi driver supply elasticities.

fee, since the latter inserts a wedge between effort and income (See, e.g., Lazear and Oyer, 2012). Our results show why it may be difficult to implement lease-type schemes in practice. While this paper focuses on the value of proportional fee schemes for rideshare drivers, our results are relevant for any job where the “right to work” can be purchased at either a flat rate or by giving up a share of earnings. For example, service professionals like hair stylists and cosmetologists face this sort of choice, working on commission or renting a salon chair. Many franchise contracts also reflect this trade-off: potential franchisees often pay a fixed cost to the franchise owner, as well as or instead of a royalty quoted as a percentage of sales.

The next section outlines a theoretical framework that contrasts incentives and constraints under the taxi and rideshare compensation schemes. Section 3 describes our experimental design and context. Section 4 presents estimates of drivers’ labor supply elasticities. Section 5 analyzes the Taxi take-up decision and shows that low take-up is best explained by loss aversion. Section 6 discusses estimates of compensating variation, comparing rideshare and leasing. Section 7 concludes.

2 Theoretical Framework

Our experiment is motivated by a stylized contrast between the compensation schemes embedded in rideshare and traditional taxi work arrangements. In Boston, until recently, Uber retained a flat fee of 20% or 25% of its drivers gross fares (referred to here as the “farebox”; these are base fares received plus any increase due to Uber’s surge multiplier; drivers who started before September 2015 were grandfathered into the lower fee). Most taxi drivers must lease a medallion (the legal right to drive) per shift, day, or week, but can then drive commission-free. Expenses (mostly gas) are paid by drivers under both schemes. Taxi medallion leases may or may not cover use of a vehicle. Uber also offered its drivers the opportunity to rent or lease cars through a program known as “vehicle solutions,” though few drivers did this.³

³Lyft offers its drivers a similar compensation arrangement. Uber changed its pay policy in June 2017 to loosen the link between rider fares and driver earnings, an innovation known as “up front pricing.” Lyft has experimented with similar schemes. Neither rideshare platform requires drivers to make advance payments analogous to medallion leasing, though some rideshare upstarts such as Fasten have tried such schemes.

2.1 Budget Sets

We capitalize “Taxi” when referring to the lease-based compensation schemes offered to Uber drivers in our experiment. This is cast against a simplified but realistic characterization of the “Rideshare” contract facing Uber drivers. Fares are cast in terms of average hourly earnings, w , taken to be the same for Rideshare and Taxi drivers. This is unrestrictive because differences in wages can be modeled as part of the Rideshare fee, or reflected in a negative fee for Taxi drivers.

Drivers drive for h hours, so their weekly farebox is wh . Their compensation schemes are as follows:

- Rideshare drivers earn $y_0 = w(1 - t_0)h$, where t_0 is the Rideshare fee.
- Taxi drivers earn $y_1 = w(1 - t_1)h - L$, where L is a Taxi lease price and $t_1 \leq 0$ reflects a possibly higher Taxi wage.

Drivers can choose not to work and earn nothing, but leases must be purchased in advance. The quantity $t_0 - t_1$ is the difference in “tax rates” implicit in the two contracts.

Our experiment ran for one week at a time, and many drivers indeed lease weekly, so it’s natural to think of L as a weekly lease, with drivers choosing Rideshare and Taxi week by week. Alternately, we can imagine Taxi as permanently displacing Rideshare or vice versa, in which case the relevant decision-making horizon might be longer, with L scaled accordingly. After laying out the basic framework, we briefly consider the contrast between Taxi and Rideshare in a life-cycle framework where the opportunity to choose between contracts may be transitory and future wages are uncertain.

Figure 1 sketches the Rideshare and Taxi budget sets when $w = 20$, $L = 100$, $t_1 = 0$ and $t_0 = .25$, so the difference in tax rates in this example is just the Rideshare fee (these are realistic values for wages and fees in Boston, but real-world medallion lease costs are much higher). In general, the budget lines cross where the farebox solves

$$wh = \frac{L}{t_0 - t_1} \equiv B,$$

a quantity we call the Taxi *breakeven*. This is \$400 in the figure, attained by drivers who drive at least 20 hours. Drivers who collect more than \$400 in fares come out ahead under Taxi, while drivers with a lower farebox take home more under Rideshare. Note that the indifference

curves sketched in this figure reflect increasing utility as the curves shift northwest. A driver with indifference curve u_0 prefers Rideshare, while a driver with indifference curve u_1 prefers Taxi.

Figure 1 compares a pair of drivers with fareboxes above and below breakeven. Drivers with hours above breakeven clearly benefit from Taxi. But some drivers with a below-breakeven farebox under Rideshare may respond to the higher Taxi wage by driving longer hours, thereby clearing breakeven. This scenario is sketched in Figure 2. As in Ashenfelter (1983)'s analysis of welfare program participation, we compute the theoretical take-up threshold by expanding an excess expenditure function that approximates the cash transfer required to attain a reference utility level.

The expenditure function for a generic labor supply problem is

$$e(p, w, \bar{u}) \equiv \min_{x, l} px + wl \text{ s.t. } u(x, l) = \bar{u},$$

giving the minimum spent on consumption (x) at price p and leisure (l) at price w in the effort to reach utility \bar{u} . Excess expenditure is spending minus the value of drivers' time endowment, T , that is:

$$s(w, \bar{u}) \equiv e(p, w, \bar{u}) - wT.$$

Using the fact that expenditure is minimized by compensated demand functions, x^c and l^c , we can write

$$s(w, \bar{u}) = px^c + wl^c - wT = px^c - wh^c.$$

The cash needed to reach a given utility level is the difference between consumption spending and driver earnings when these quantities are chosen optimally.

We model Rideshare and Taxi in this framework by treating lease costs and ride-hailing fees as parameters in an expanded excess expenditure function. Ignoring other earnings opportunities for the moment, the cash transfer needed to attain \bar{u} when driving under a scheme with L and t as parameters can be written

$$f(w, \bar{u}; t, L) = (px^c + L) - w(1 - t)h^c = s(w[1 - t], \bar{u}) + L.$$

Let u_0 denote utility attained when driving Rideshare, a contract described by $L = 0, t = t_0$. Drivers prefer Taxi when the Taxi contract allows them to reach u_0 for less than $f(w, u_0; t_0, 0)$.

Specifically, assuming $t_1 = 0$, Uber drivers take Taxi when

$$\underbrace{f(w, u_0; 0, L)}_{\text{Taxi}} < \underbrace{f(w, u_0; t_0, 0)}_{\text{Rideshare}},$$

or, equivalently, when

$$s(w, u_0) + L < s(w_0, u_0), \quad (1)$$

where $w_0 = w(1 - t_0)$ is the after-fee Uber wage. Using a second order expansion of $s(w, u_0)$ around $s(w_0, u_0)$ and simplifying using Shephard's Lemma, the Taxi participation rule is

$$L - t_0 w h_0 - \frac{1}{2} \left(\frac{\partial h^c}{\partial w} \frac{w(1 - t_0)}{h_0} \right) t_0 w h_0 \frac{t_0}{1 - t_0} < 0, \quad (2)$$

where h_0 is Rideshare labor supply (we omit the superscript reminding us this is the level of work determined by the compensated supply function).

It's useful to rewrite the Taxi participation rule in terms of the Taxi breakeven:

$$\underbrace{w h_0}_{\text{Rideshare farebox}} > \underbrace{\frac{L}{t_0} \left(1 + \frac{\delta}{2} \frac{t_0}{1 - t_0} \right)^{-1}}_{\text{adjusted breakeven}}, \quad (3)$$

where δ is the substitution elasticity evaluated at the after-fee Rideshare wage:

$$\delta \equiv \frac{\partial h^c}{\partial w} \frac{w(1 - t_0)}{h_0} = \frac{\partial h^c}{\partial w} \frac{w_0}{h_0}$$

and $w_0 = (1 - t_0)w$. This shows that a positive substitution elasticity reduces the participation threshold by the proportional amount

$$\frac{1}{1 + .5\delta \frac{t_0}{1 - t_0}}.$$

Eligible drivers with a Rideshare farebox that clears breakeven should always prefer Taxi. But some with a farebox below breakeven should also accept a Taxi contract. With a unit-elastic compensated response and a fee of 25%, for example, we expect the participation threshold to be reduced relative to breakeven by $1 - \frac{1}{1 + .25/2 \times .75} \approx 14.5\%$.

2.2 Compensating for Taxi-Type Compensation

To model driver choices between work arrangements, we derive the payment required to make up for loss of the opportunity to drive under a proportional fee-based contract. This is compensating variation (CV), where the baseline condition is the Rideshare budget line with

an interior solution and the alternative is the Taxi budget set. Positive CV means payment is required for imposition of Taxi, while negative values arise for drivers who prefer Taxi. Although CV is tied to the specifics of the compensation scheme on offer, the results of our experimental Taxi-Rideshare comparisons can be used to extrapolate compensation values to other markets where workers might choose between paying a proportional tax on their earnings and paying a fixed up-front fee.

Formally, CV is the difference in cash required to reach a reference utility level given the Taxi and Rideshare budget lines:

$$f(w, u_0; 0, L) - f(w, u_0; t_0, 0),$$

where u_0 is the Rideshare utility level. Using Shephard's Lemma as in equation (2), the CV required as compensation for Taxi can be shown to be

$$CV = \{L - t_0wh_0\} - t_0wh_0 \frac{\delta t_0}{2(1 - t_0)}. \quad (4)$$

Rideshare drivers for whom CV is negative take the Taxi scheme when offered, producing the participation rule described by (2).

A Leontief ($\delta = 0$) driver should be paid the difference between his or her lease costs and Rideshare fees. Elastic labor supply favors Taxi, reducing CV. Even so, the principal determinant of CV for most drivers is likely to be $L - t_0wh_0$, the difference between lease costs and Rideshare fees. This difference is largest for Uber and Lyft's many low-hours drivers. Recall also that in the absence of substantial income effects on the demand for leisure, CV approximates the difference in driver surplus yielded by the two compensation schemes (this in turn equals the corresponding equivalent variation).

The left panel of Figure 3 illustrates the CV calculation generated by a move from the Rideshare to Taxi budget lines. A Rideshare driver working at point A drives 10 hours and is on indifference curve u_0 . Faced with a Taxi budget line, this driver drives 13 hours, but is worse off on u_1 . It seems natural to compensate this driver by an amount equal to the excess of his lease over what he used to pay in Rideshare fees. But a payment of $L - t_0wh_0$ puts non-Leontief drivers above point C on u_0 , as indicated by the blue line extending from point A with a slope equal to the Taxi wage. Payments equal to lease costs minus *ex ante* Rideshare fees over-compensate for Taxi because the Taxi scheme increases wages, yielding

additional driver surplus. The term $wh_0 \frac{\delta t_0}{2(1-t_0)}$ in equation (4) captures this surplus, a term denoted by σ in Figure 3. The surplus generated by higher Taxi wages is the product of the proportional Taxi wage advantage, $\frac{t_0}{1-t_0}$, the substitution elasticity (δ), and driver fees, t_0wh_0 . This product approximates the area under the driver's supply curve between his net-of-fee Rideshare and Taxi wages.

Choosing Not to Drive

The compensation formula above presumes Rideshare drivers accept the Taxi budget line as a condition for compensation. But we might instead allow former Rideshare drivers to refuse Taxi, taking some of their compensation in the form of increased leisure. In this scenario, drivers are made whole by “unemployment insurance” (UI) in an amount that takes them to the u_0 ordinate, a scenario illustrated in the right panel of Figure 3.

To compute the compensation needed in this case, we assume the marginal utility of leisure is zero at $h = 0$, so drivers with a wage of zero choose zero hours. Expanding the excess expenditure function for Rideshare utility with a wage of zero around Rideshare expenditure with a fee of t_0 , we have:

$$s(0, u_0) = s(w_0, u_0) + (-h_0)(-w(1-t_0)) - \frac{1}{2} \frac{\partial h^c}{\partial w} w^2 (1-t_0)^2. \quad (5)$$

By definition of u_0 , Rideshare drivers with no unearned income and no lease to cover have consumption equal to their Rideshare earnings, so $s(w_0, u_0) = 0$. The compensation required for the replacement of Rideshare work opportunities with UI is therefore

$$UI = (1-t_0)wh_0 - \frac{1}{2} \left(\frac{\partial h^c}{\partial w} \frac{w(1-t_0)}{h_0} \right) ([1-t_0]wh_0) \quad (6)$$

$$= (1-t_0)wh_0 \left[1 - \frac{\delta}{2} \right] \quad (7)$$

The replacement rate for lost Rideshare earnings in this case is approximately one minus half the compensated labor supply elasticity. For Leontief drivers, the replacement rate is 100% since their $\delta = 0$.

2.3 Life Cycle Considerations

We compare Rideshare and Taxi in a multi-period setting using the Browning et al. (1985) duality framework built around the profit function. Just as the excess expenditure function is the potential function for compensated labor supply at a fixed utility level, the profit function is the potential function for Frisch labor supply. These supply functions characterize the response to perfectly anticipated wage changes (MaCurdy (1981) calls these “evolutionary” wage changes) or to transitory changes that have little effect on lifetime wealth (more precisely, little effect on the marginal utility of lifetime wealth). The derivative of Frisch labor supply with respect to the wage rate is the Frisch or intertemporal substitution elasticity (ISE).

With intertemporally additive preferences and a known path for wages, workers’ total profit functions are given by the sum of period- s profit functions, $\pi_s(r, w_s, p_s)$, defined as

$$\pi_s(r, w_s, p_s) \equiv \max_{u, x, l} ru + w_s(T - l) - p_s x; \quad u = v_s(x, l),$$

where r is the reciprocal of the marginal utility of wealth, $v_s(x, l)$ is period s utility, and wages and prices in period s are time-varying. The profit function imagines consumers valuing their utility at price r ; profit is then the monetary value of utility plus earnings, net of expenditure on inputs in the form of consumption.

Consider a driver making a life-cycle plan in the face of known wages and prices, choosing between Rideshare and Taxi at time (week) s . This driver prefers Taxi if the Taxi contract is profitable for that week. That is, Taxi beats Rideshare in week s when

$$\pi_s(r, w_s) - \pi_s(r, w_s[1 - t_0]) > L.$$

This comparison presumes the utility price is unchanged by Taxi, either because the Taxi opportunity and parameters are known at the time plans are made, or because the Taxi option is short-lived. We assume goods prices are constant, so p_s is left in the background.⁴

Expanding $\pi_s(r, w_s)$ around the value of Rideshare profits, $\pi_s(r, w_s[1 - t_0])$, the life-cycle

⁴Our streamlined notation also ignores the the fact that wage and price variables determining profits in a future period s are discounted back to the decision-making date; see Browning, Deaton, and Irish (1985) for details.

participation rule for Taxi at week s is approximated by

$$\frac{\partial \pi_s(r, w_s[1 - t_0])}{\partial w} w_s t_0 + \frac{1}{2} \frac{\partial^2 \pi_s(r, w_s[1 - t_0])}{\partial w^2} (w_s t_0)^2 > L. \quad (8)$$

Applying a life-cycle version of Shephard's lemma, this can be written

$$\underbrace{w_s h_{s0}}_{\text{Rideshare earnings}} > \underbrace{\frac{L}{t_0} \left(1 + \frac{\delta^f}{2} \frac{t_0}{(1 - t_0)} \right)^{-1}}_{\text{life cycle breakeven}}, \quad (9)$$

where $\delta^f \equiv \frac{\partial h_s^f(r, w_s)}{\partial w_s} \frac{w_s(1-t_0)}{h}$ and $h_{s0} \equiv h_s^f(r, w_s[1 - t_0])$ is Frisch labor supply for Rideshare in period s . The earlier Taxi participation rule therefore stands, but with the Hicks substitution elasticity replaced by the possibly larger ISE, δ^f .

The revision to CV in a life-cycle framework parallels that for participation. Specifically, CV is the sum of the difference in within-period profits:

$$CV = [\pi_s(r, w_s) - L] - \pi_s(r, w_s[1 - t_0]).$$

Using the expansion yielding equation (8), this becomes:

$$CV = \{L - t_0 w_s h_{s0}\} - t_0 w_s h_{s0} \frac{\delta^f t_0}{2(1 - t_0)}. \quad (10)$$

This is the same as (4), with the ISE δ^f again replacing the substitution elasticity, δ . Since the ISE (weakly) exceeds the Hicks substitution elasticity, a life-cycle perspective tends to favor Taxi. Because our experimental design offers temporary wage changes, we interpret the experiment as identifying δ^f .

In practice, drivers considering a weekly lease must do so without knowing next week's wage or farebox. Suppose that a Rideshare driver who doesn't know next week's wages is offered the opportunity to buy a one-week lease. Although marginal utility of lifetime wealth presumably changes little as a result of wage surprises, some idea of w is required to make a wise near-term choice. Knowing how much a driver drives in response to each wage, a predicted wage implies a predicted farebox. The econometric framework outlined in Section 5, below, therefore embeds farebox prediction in an empirical model for Taxi participation.

2.4 Outside Options

The drivers in our experiment can typically supply as many hours to Uber as they like at the implicit market wage, but many Rideshare drivers work at another job (Hall and Krueger, 2017). We model movement between a single Rideshare employer and alternative employment as motivated by declining earnings opportunities on the alternative job. For alternative jobs with institutional limits on hours, such as shift work or salaried office work, the decline is likely to be precipitous. On other sorts of jobs, including alternative ride-hailing platforms, any pay advantage over Uber may taper smoothly. We might imagine, for example, that Lyft takes lower fees than Uber, but offers its drivers less steady trip demand. This market structure is captured by assuming that drivers earn $e(a)$ for a hours worked on an alternative job, where $e(a)$ is increasing but concave.⁵

The excess expenditure function for a driver who holds an alternative job is

$$s^a(p, w, \bar{u}) = \min_{x, h, d} px - wh - e(a) \text{ s.t. } u(x, T - h - a) = \bar{u},$$

where the a superscript indicates that this is excess expenditure for someone who works an alternative job. As always, excess expenditure is minimized by the compensated demand functions x^c, h^c, a^c , so

$$s^a(p, w, \bar{u}) = px^c - wh^c - e(a^c).$$

Writing $f^a(w, \bar{u}, L, t)$ for the cash required to reach utility \bar{u} in this scenario, when faced with driving (Rideshare or Taxi) wage w and parameters L, t , yields the relevant excess expenditure functions:

- Rideshare: $f^a(w, \bar{u}; t_0, 0) = px^c - w(1 - t)h^c - e(a^c) = s^a(w(1 - t_0), \bar{u}) = s^a(w_0, \bar{u})$
- Taxi: $f^a(w, \bar{u}; 0, L) = (px^c + L) - wh^c - e(a^c) = s^a(w, \bar{u}) + L,$

where it's understood that compensated demand functions are different in the two schemes. The appendix derives the usual Shephard's lemma result in this context:

$$\frac{\partial f^a}{\partial w} = \frac{\partial s^a}{\partial w} = -h^c, \tag{11}$$

⁵This setup is inspired by the Gronau (1977) model of home production, where workers get utility from a single consumption good and from leisure, and can produce the consumption good under diminishing returns at home or buy it with money earned on a job paying constant wages.

with the modification that compensated labor supply now includes only hours worked as a driver.

We can also use Shephard’s lemma to show that Rideshare drivers with alternative jobs are happy to drive Taxi when:

$$wh_0 > \frac{L}{t_0} \left(1 + \frac{1}{2(1-t_0)} \tilde{\delta} t_0 \right)^{-1}. \quad (12)$$

This looks like (3), but the substitution elasticity in this case, denoted by $\tilde{\delta}$, measures the change in *hours driving for Uber*, while total labor supply includes hours driving plus hours worked on the alternative job, $H = h + a$. The formula for CV is adjusted similarly. The wage elasticity of hours driving for a particular platform is likely to be larger than the elasticity of total hours worked since changes in H may reflect substitution from h to a with little change in H . Because our experiment measures the change in hours driving for Uber, this change in the interpretation of parameters leaves our welfare analyses unchanged.

3 Experimental Design

Uber and its ride-hailing competitors routinely offer drivers temporary increases in pay (known as “promotions”) that are designed to boost the supply of trips. We estimate labor supply elasticities and lease aversion parameters using a randomized experiment presented to drivers as an Uber promotion called the *Earnings Accelerator*.

3.1 Overview

The experiment unfolded in three phrases: (1) selection of eligible drivers, (2) opt-in treatment weeks, and (3) Taxi treatment weeks. The opt-in phase was designed to select drivers who seem likely to take note of Uber’s promotional messaging and to obtain informed consent for exposure to Taxi offers. Initial ISE estimates from the opt-in phase also allowed us to calibrate lease prices. Finally, the Taxi phase identified drivers willingness to lease and provided further evidence on intertemporal substitution. Online Appendix Table A1 sketches the experimental timeline.

Drivers were eligible for inclusion in the experiment if they took at least four trips and drove an average of 5-25 hours per week in the four weeks prior to selection (the last 3

weeks of July and the first week of August 2016). The omission of higher hours drivers—those with average weekly hours above 25—reduced experimental costs and allowed us to focus on a sample of drivers with farebox values clustered around modest Taxi breakevens. Higher hours drivers may differ from other drivers, but an analysis of drivers grouped by hours driven within the eligible sample shows little systematic variation in the behavioral parameters that contribute to the computation of CV.

Roughly 45% percent of Boston drivers were eligible for inclusion in the experiment. Although the cap on hours per week reduces average hours in the eligible sample relative to the city average, drivers in the eligible sample are otherwise similar to the pool of active Boston drivers (that is, the group who took at least four trips in the previous month). For example, 14% of both the active and eligible samples are female and both groups had used the Uber platform for an average of 14 months. These comparisons appears in the first two columns of Table 1.

A total of 1600 eligible drivers were selected for inclusion in the experiment. The experimental design randomized within strata defined by average hours driven in July 2016, driver fee class (commission rate), and vehicle model year. The low hours stratum includes drivers who averaged 5-15 hours per week, while the high hours group averaged 15-25 hours per week. The 20% fee class includes veteran drivers who signed up before September 2015, while other drivers paid a commission rate of 25%. Because Lyft requires its drivers to use cars no older than 2004, our strata distinguish between drivers with cars from model year 2003 or older and drivers with newer Lyft-eligible cars. We also report the proportion of drivers with cars newer than 2010 since Lyft’s most important promotion requires drivers operate newer vehicles. Drivers were randomly sampled and randomly assigned to the first or second opt-in week within these three strata. As can be seen in column 4 of Table 1, which reports strata-adjusted differences in means, the experimental sample has characteristics similar to those of drivers in the rest of the eligible sample.

3.2 Opt-In Weeks

The 1600 sampled drivers were offered the opportunity to drive for one opt-in week with no Uber fee. Half of the drivers (Wave 1) were offered a week of fee-free driving in the first opt-in week. As the first wave was driving fee-free, drivers the second half sample (Wave 2)

were offered the opportunity to opt in for fee-free driving the following week. This design was meant to control for wealth effects induced by the higher fee-free wage. Appendix Table A2 shows that driver characteristics are well balanced across waves.

Drivers in both waves were offered fee-free driving by e-mail, text message and in-app notification on Monday morning of the relevant opt-in week; they had until midnight the following Saturday to opt-in. Sampled drivers received up to three emailed reminders to opt-in by the deadline. Drivers who opted in paid no Uber fee on all trips taken in the subsequent week. This was reflected in their immediate in-app trip receipts and weekly pay statements (participating drivers saw a fee of zero in receipts and statements). Fee-free driving increased a driver’s total payout by 25% in the 20% fee class ($.25 = \frac{1}{.8} - 1$) and by 33% in the 25% fee class ($.33 = \frac{1}{.75} - 1$).

Opt-in statistics for each of the free week strata are reported in Table 2. Roughly sixty-four percent (1031/1600) accepted the opportunity to drive fee-free. Although fee-free driving should be attractive to all drivers, many appear to ignore Uber messaging beyond the offer of trips. This likely reflects the fact that (during our experimental period) Uber drivers received many electronic messages each week – some of which they undoubtedly ignored.⁶ The need for intense messaging to garner driver attention is reflected in a decline in take-up from Wave 1 (71%) to Wave 2 (58%), after we stopped the opt-in reminders mid-week. Messaging was reduced in view of higher-than-expected take-up and a consequent risk of running over-budget. Discussions with Uber’s Boston team suggest our take-up rates compare favorably with the response rate to other no-lose driver promotions requiring an opt-in. Incomplete take-up may also reflect the fact that drivers who opted in consented for their data to be used in academic research and to receive further Earnings Accelerator offers.

Table 3 shows that drivers who opted in drove and earned more than other drivers during the opt-in week. In the pooled sample including both high and low hours drivers, those who opted in had an opt-in-week farebox roughly \$100 higher than the farebox of drivers who opted out. Those who opted-in also drove 4 more hours that week. On the other hand, these gaps are much smaller when averaged over the month of July. This is consistent with the idea that inattention drove low opt-in rates: drivers who drove during opt-in week had more chances to see the offer through the in-app notification. In addition to higher average

⁶In view of this, Uber moved later to cap the number of promotion-related messages sent to drivers.

hours in the opt-in group, participating drivers are a little younger. Other characteristics, including average commission rates, percent female, and months on platform, are unrelated to participation status.

3.3 Taxi Treatments

The Taxi phase of the experiment offered random subsets of the 1031 drivers who opted in to fee-free driving the opportunity to buy additional weeks of fee-free driving for a modest lease. These Taxi treatments were randomly assigned within strata defined by average hours and fee. Eight treatments were offered in each Taxi week, two for each hours/fee combination. Online Appendix Tables A3 and A4 show that random assignment of these treatments balanced the characteristics of drivers in the Taxi treatment and control groups.

Each taxi treatment consists of a fee reduction, $t_1 - t_0$, and a lease price, L . Based on the ISE estimates computed using data from opt-in weeks, lease prices were calibrated so as to be attractive to roughly 60% of drivers in each stratum. These lease prices and fee changes are listed in Table 4. In the first Taxi week, 40% of drivers in each stratum were offered the opportunity to buy another week of fee-free driving and 20% were offered negative fee driving in the form of a 12.5% wage increase ($t_1 = -.125$). Lease prices in the first Taxi week ranged from \$45 to \$165. The treatments in week 2 were less generous—the negative fee treatment was replaced with a half fee treatment—but also less expensive, with leases priced between \$15 and \$60. Each week 2 treatment was offered to 30% of drivers within strata. Figure 4 summarizes experimental staging and design parameters.

As during with solicitation for fee-free driving in opt-in week, drivers were offered Taxi contracts via e-mail, text messages and in-app notifications. These offers were sent one week in advance and highlighted the breakeven amount. For example, drivers in the 25% fee class who were offered a half-fee treatment for \$35 were told “As long as your weekly total fares+surge exceed \$280, you’ll come out ahead.” Email and text messages included links that clicked through to a simple table contrasting the former and the revised fee calculation for a sample trip. Emails and text messages also included links that clicked through to a calculator that showed net earnings with and without the treatment for any driver-selected value of fares plus surge. Figure A1 in the Online Appendix shows a message delivered in the Taxi promotions; Figure A2 shows the calculator.

Drivers who opted-in to taxi had lease payments deducted from their pay for the opt-in week. This deduction appeared as a negative entry on otherwise-standard weekly pay statements on the line that shows any (usually positive) payment drivers might have earned through Uber promotions.⁷ These deductions were labeled “Earnings Accelerator buy-in.” During the following week, when taxi was “live” for those who opted in, drivers were able to see the benefits of fee reductions in real time; each trip receipt that week reflected the reduced fee (see Figure 5).

It is worth noting that our experimental lease amounts are well below the price of a traditional taxi medallion lease: before the advent of ride-hailing, Boston medallion leases (including vehicle) ran around \$700/week and over \$100/day. Our virtual medallions were priced from \$50-\$165/week. These prices were calibrated to appeal to drivers with weekly earnings in particular ranges, as explained below. As a measure of the contemporary empirical relevance of our design, it’s noteworthy that in 2016 a Boston ride-hailing upstart (Fasten) offered its drivers the option to pay \$80/week or \$15/day to drive fee-free.⁸

4 Labor Supply Effects

4.1 Impacts on Participants

Our estimates of the ISE rely on a linear labor supply model for drivers with positive hours driven. We preface these estimates with a set of results using offers as instruments in a two-stage least squares (2SLS) setup that captures the impact of Earnings Accelerator *participation* on measures of labor supply. Experimental participation is defined in two ways.

⁷Drivers who earned less than needed during opt-in week carried a negative balance into the following pay period. This was an unusual scenario since leases were inexpensive and most drivers who opted in drove during opt-in week.

⁸In 2010, the Boston medallion lease cost for a single driver was capped at \$700/week, \$139/day, and \$77/12-hour shift (BPD Circular Date 12-30-09 “2010 Standard Shift Rental Agreement”). Newer cars leased for an additional 170/week. Drivers could split a weekly lease for no more than \$800. Before the advent of ride-hailing, short supply meant medallions typically leased at the cap. Side payments to Boston fleet owners also appear to have been entrenched (See the 2013 Boston Globe stories linked under <http://www.bostonglobe.com/metro/specials/taxi>). Data on medallion prices is spotty; a CommonWealth Magazine article (<http://commonwealthmagazine.org/transportation/taxi-medallion-owners-under-water-and-drowning/>) quotes a pre-ride-hailing Boston medallion price of over \$700,000, down recently to about half that. NYC medallion prices are said to have peaked at over one million dollars (<http://seekingalpha.com/article/3177766-taxi-farebox-declines-a-harder-hit-to-medallion-owner-bottom-lines?page=2>).

First, using the full research sample of 1600, participants are drivers who agreed to be in the experiment during the opt-in phase. Second, for the 1031 drivers who opted in, participants are those who purchased a Taxi contract. Participation estimates distinguish extensive from intensive margin effects, identify possible changes in average hourly compensation, and reveal any possible anticipatory or post-treatment labor supply changes that might signal confounding wealth effects. The participation analysis yields three important findings: (1) Earnings Accelerator participation had no effect on the extensive margin (that is, effects on whether drivers drive at all); (2) Participation boosted hours driven and driver earnings considerably during treatment weeks, increases that are proportional, with no corresponding change in average hourly earnings; (3) We see no evidence of anticipatory or post-experiment effects in the treated group.

The analysis sample for 2SLS estimation of participation effects stacks data for two pairs of weeks: the first pair contains data on 1600 drivers from the first two waves; the second pair includes observations from the two Taxi weeks for the 1031 drivers who opted in to fee-free driving and agreed to receive Taxi offers later. The endogenous variable in this setup, D_{it} , indicates fee-free driving in week t or purchase of a Taxi contract during the Taxi opt-in weeks, to be used in week t . The instrument, Z_{it} , indicates offers of fee-free driving or a Taxi contract in week t . For example, Z_{i1} is switched on for the 800 drivers offered fee-free driving in Wave 1 and for the 619 drivers offered a Taxi lease during the first week of the Taxi trial.

For a set of weekly labor supply outcomes denoted by Y_{it} , the 2SLS setup for ISE estimates can be written:

$$Y_{it} = \alpha D_{it} + \beta X_{it} + \eta_{it} \quad (13)$$

$$D_{it} = \gamma Z_{it} + \lambda X_{it} + v_{it}, \quad (14)$$

where X_{it} includes dummies indicating the strata used for random assignment, driver gender, the number of months a driver had worked on the Uber platform, one lag of log earnings, and indicators for whether a driver used Uber’s “vehicle solutions” leasing program and whether a driver had a car from model year 2003 or older. Because drivers appear in the sample more than once, standard errors in this setup are clustered on driver.

As can be seen in Figure 7, fee-free driving and Taxi offers both boosted participating

drivers' hours and farebox considerably, with little effect on the extensive margin (that is, on an indicator for any Uber activity, $wh_{it} > 0$). The upper panel of the figure also suggest that fee-free driving had no effect on participants' hours, farebox, and Uber activity rates in the week before Wave 1 (this is opt-in week for Wave 1) or in the week following fee-free driving for Wave 2.⁹ In weeks of fee-free driving, however, participating drivers' hours and farebox rose by about 35%, though their extensive margin activity rates were almost unchanged. The estimates behind Figure 7, reported in Panel A of appendix Table A5 show an effect of .04 on Uber activity during opt-in week. The absence of an effect before and after treatment weeks weighs against significant wealth effects from higher wages during one of the two treatment weeks. This is not surprising given that the earnings gains are small, relative to lifetime wealth.

The lower panel of Figure 7 shows that the Taxi treatment had a similar, though slightly smaller, effect on hours and farebox of around 30%. Effects on hours and earnings were smaller in the 2nd week of Taxi than in the first, most likely reflecting the fact that the treatments offered that week were less generous. Appendix Table A6 shows that 2SLS estimates of participation effects are reasonably similar across hours groups. For example, the more precisely estimated effects in models with covariates show increases of .43 and .34 in the high and low hours groups in response to Taxi participation and .30 and .34 in the high and low groups during opt-in weeks. The estimated effect of Taxi participation on Uber activity is .01, and not significantly different from zero.

The fact that the farebox and earnings effects plotted in Figure 7 are similar suggests that Uber drivers face fairly constant rideshare earnings opportunities, as hypothesized in Section 2.4. Appendix Table A6 reports 2SLS estimates of Earnings Accelerator participation effects on average hourly farebox and other measures of driver effort and labor supply, including the number of completed trips, the number of days worked during the week, the proportion of weekly trips earning a surge premium, and the average rating on rated trips during the week. Consistent with the hours and earnings estimates, these results show clear experimental effects on completed trips and the number of days with any driving. Effects on other outcomes, however, including average hourly farebox and ratings, are small and not

⁹While the Euler equation for hours would predict *some* reduction in labor supply in the week before and in weeks after the wage increase, this would be too small to detect in our setting.

significantly different from zero. Appendix Figure A3 reports participation effects on the distribution of hours worked.

4.2 Estimating the Rideshare ISE

The ISE for Rideshare hours is estimated by replacing the dependent variable in (13) with log hours driving, and replacing the endogenous variable in (14) with log wages earned as a driver. The hours variable, h_{it} , measures weekly hours with the Uber app toggled on; the wage, w_{it} , is the average hourly farebox net of Uber fees. The 2SLS estimate of the coefficient on $\ln w_{it}$ is our measure of the ISE, denoted δ^f in Section 2 (this is the parameter $\tilde{\delta}$ in the model with alternative jobs). Life-cycle logic suggests wealth effects from leasing should be small, so offers of Taxi leasing and fee-free driving should generate similar ISEs when estimated in the same population.

The first stage effect of Earnings Accelerator offers on log wages (γ in equation 14) depends on: (1) the experimental participation rate, and (2) the magnitude of experimentally-induced fee changes. To see this, let w_{it}^0 denote a driver's potential average hourly farebox in the absence of treatment. Participation decisions determine average hourly earnings through

$$\begin{aligned} w_{it} &= w_{it}^0(1 - t_0)(1 - D_{it}) + w_{it}^0(1 - t_1)D_{it} \\ &= w_{it}^0(1 - t_0) + w_{it}^0(t_0 - t_1)D_{it}. \end{aligned}$$

Ignoring covariates and using the fact that randomly assigned treatment offers are independent of w_{it}^0 , the first stage effect of offers on wages is

$$\begin{aligned} E[w_{it}|Z_{it} = 1, t_0, t_1] - E[w_{it}|Z_{it} = 0, t_0, t_1] \\ = (t_0 - t_1)E[w_{it}^0|D_{it} = 1] \times P[D_{it} = 1|Z_{it} = 1]. \end{aligned} \tag{15}$$

In other words, wages go up in the treatment group in an amount given by the experimental fee change times average wages for participants times the opt-in rate.¹⁰

The experimentally-induced proportional change in wages is obtained by dividing (15) by average hourly earnings for controls, $E[w_{it}|Z_{it} = 0] = E[w_{it}^0](1 - t_0)$. Assuming wages are similar for participants and other drivers, a claim supported by Table 3, the proportional

¹⁰The derivation here uses the fact that $D_{it} = 1$ implies $Z_{it} = 1$, which in turn yields $E[w_{it}^0|D_{it} = 1, Z_{it} = 1] = E[w_{it}^0|D_{it} = 1]$.

wage increase generated by the Earnings Accelerator is:

$$\frac{E[w_{it}|Z_{it} = 1, t_0, t_1] - E[w_{it}|Z_{it} = 0, t_0, t_1]}{E[w_{it}|Z_{it} = 0, t_0, t_1]} = \frac{(t_0 - t_1)}{1 - t_0} P[D_{it} = 1|Z_{it} = 1]. \quad (16)$$

In other words, the proportional first stage for wages is the experimentally-induced change in fee divided by the baseline take-home rate, times the treatment take-up rate. For example, with a take-up rate of $2/3$, the proportional first stage for an experiment that eliminates a 25% fee is roughly $\frac{.25}{.75}.66 = .22$.¹¹

Equation (16) is the first stage for a just-identified 2SLS estimator using a single dummy instrument. Over-identified estimates using multiple instruments that distinguish different sorts of offers and different experimental weeks should generate more precise estimates. Fee-free driving offers were made twice, once in each opt-in week, providing a pair of instruments to identify the ISE using data from opt-in weeks. Taxi offers produce 16 instruments, one for each lease, tax rate, and hours stratum in each of two Taxi weeks. We compute just-identified and over-identified estimates of the ISE in models controlling for random assignment strata and for a set of driver covariates listed in table notes. A parallel set of 2SLS estimates controlling only for strata appears in the appendix.

Just-identified estimates of the ISE range from about 1.2 using data from opt-in week to 1.8 in the Taxi sample. These estimates, reported in Panel A of Table 5, are not too far from the experimentally-identified ISE estimates reported for Swiss bicycle messengers by Fehr and Goette (2007).¹² The over-identified estimate of the (pooled-sample) ISE using Taxi variation falls to about 1.4, still larger than the corresponding estimate using the full sample. It's perhaps unsurprising that drivers who find Taxi leasing attractive are more elastic.¹³ In both samples, the just-identified and over-identified estimates are precise enough to rule out much smaller values. Moreover, we see little in the way of systematic elasticity differences between low and high hours drivers. It's also noteworthy that the corresponding

¹¹The first stage in logs is $\ln \frac{1-t_1}{1-t_0} \times P[D_{it} = 1|Z_{it} = 1]$, but $\ln \frac{1-t_1}{1-t_0} \approx \frac{(t_0-t_1)}{1-t_0}$.

¹²Fehr and Goette (2007) estimate an ISE of between 1.12 and 1.25 for an all-male sample that is probably younger than our sample of Uber drivers.

¹³The argument that leads us to expect more elastic Taxi drivers parallels the phenomenon of selection on moral hazard in health insurance markets. Einav, Finkelstein, Ryan, Schrimpf and Cullen (2013) argue that health insurance plans are chosen partly in view of anticipated healthcare utilization while covered by insurance.

OLS estimates of equation (13), reported in Panel B, are far smaller than the ISEs identified by random assignment.¹⁴

Two further comments on the impressively elastic behavior of Boston Uber drivers are in order. First, the ISE estimation sample omits drivers with no hours in a given week. Because Earnings Accelerator offers are largely unrelated to drivers' decisions as to whether to be active at all (a result shown in Figure 7), this extensive-margin conditioning seems innocuous.

Second, as discussed in Section 2.4, the increase in Uber effort may come at the expense of work hours supplied elsewhere. Job-shifting to take advantage of higher Uber wages leaves our welfare analysis unchanged (the relevant substitution elasticity reflects changes in Uber hours). But shifting in response to higher Uber pay means our estimates of the ISE can be expected to be larger than those estimated using data on total hours worked. The most elastic alternative job response is likely to be reduced hours driving for Lyft. An appendix on platform substitution therefore reports estimates for drivers less likely to shift away from Lyft. These estimates differ little from those discussed in the text.

5 Taxi Participation

Figure 8 plots observed Taxi participation rates against predicted take-up for each of the sixteen Taxi contracts (four hours strata and commission groups times two treatments per group, in each of two weeks) offered to the sample of 1031 drivers who opted in. Predicted participation is calculated using (12), with the pre-experiment opt-in week farebox playing the role of wh_0 . A value of $\delta^f = 1.8$, taken from column 4 in Table 5, is used to compute the driver surplus produced by higher Taxi wages.¹⁵ The regression of observed participation rates on

¹⁴While we uncover little heterogeneity in the ISE by hours worked, our experiment includes drivers with a limited range of usual hours worked. One potential explanation for why our estimates of the ISE are smaller than previous estimates for taxi drivers is that, by design, we exclude full-time drivers (Farber, 2015).

¹⁵This is the estimate for drivers participating in the experiment (that is, they agreed to receive Taxi offers). Figure 8 uses fareboxes of control drivers in the same hours stratum and commission as treated drivers. The predicted participation rate for a treatment characterized by $[L, t]$ is

$$\frac{1}{N_j} \sum_{i=1}^{N_j} 1 \left\{ \log wh_{0i} > \log \left[\frac{L}{t} \left(1 + \frac{1}{2} \delta^f \frac{t}{1-t} \right)^{-1} \right] \right\}$$

where wh_{0i} is opt-in week farebox for driver i in hours/commission group j , and N_j is the size of the group. This rate therefore conditions on positive hours during opt-in week.

predicted participation rates plotted in Figure 8 shows that empirical Taxi participation rates average well below predicted. Note that this is true even for drivers in the 20% commission group (red points), who had been on the platform for over a year. It is also true for high hours drivers (diamond points).

Perhaps the drivers who skipped Taxi did so because they correctly anticipated little benefit from a Taxi contract. This possibility is explored in Table 6, which reports average earnings gains for drivers who did and did not buy a Taxi lease. The sample here is limited to the 1031 drivers who initially opted-in. Columns 1-2 use the opt-in week earnings distribution to compute the earnings gains drivers could have expected under Taxi. Expected gains are computed using an ISE of 1.2, the estimate for Earnings Accelerator participants (this adjustment is minor). For example, column 1 shows that 78 percent of drivers who accepted a Taxi contract would have expected to gain if they used opt-in week earnings to evaluate Taxi. This proportion is lower for those who did not buy a Taxi contract – 56 percent in column 2 – but still substantial. Among those expecting gains, the average gain amounts to \$92 for Taxi participants and \$66 for the non-participating group.

It’s noteworthy that Taxi participation gains forecast based on driving behavior during opt-in week are similar to those computed using the realized Taxi week earnings distribution. This can be seen by comparing the gains estimates in columns 1 and 2 of Table 6 with the estimates in columns 3 and 4 (column 3 uses realized gains for participants; column 4 is the expected gain for non-participants). However, as can be seen in panel B of Table 6, conditional on driving (most drivers in the sample indeed drove, with or without Taxi), the expected gains from a Taxi contract among non-participants were a little *larger* than the gains anticipated or realized by participants. Compare, for example, 103 and 106 dollars in anticipated benefits when forecast using the opt-in week distribution and 97 dollars gained for participants and 115 dollars in expected gains foregone for non-participants using Taxi week data.¹⁶

¹⁶Otsuka and Murakami (1987) find that consumption commitments may explain the use of share-cropping among Bangladeshi workers and farmers. Our design mitigates the importance of such commitments since most drivers participating in the first Taxi trial received higher-than-normal paychecks during opt-in week. Moreover, Online Appendix Table A8 shows that the lease aversion parameters are similar in both Taxi weeks..

5.1 Risk Aversion and Lease Aversion

Drivers’ private information does not appear to account for low Taxi take-up. Perhaps risk aversion explains why so many drivers passed up a profitable opportunity to reduce their Rideshare fees in return for a modest payment . Risk aversion seems a natural hypothesis since fee elimination increases the proportional variance of earnings by $\frac{1}{(1-t_0)^2}$. Rabin (2000) shows, however, that globally concave utility is unlikely to produce a coherent account of choices over small gambles like the one induced by our experiment (Chetty 2006 extends this argument to labor supply).

The Appendix uses data on expected gains and week-to-week farebox variation to calibrate the coefficient of relative risk aversion needed to explain low take-up among drivers for whom the expected gain from Taxi participation was positive. As in Sydnor (2010)’s investigation of homeowners’ choice of insurance deductibles, our calibration suggests drivers must be implausibly risk averse for concave utility alone to explain Taxi undersubscription.

On the other hand, loss aversion is a compelling explanation of low Taxi take-up: lease purchase looks to be a gamble that drivers hate to lose. The Appendix sketches a simple model of loss aversion in the spirit of Fehr and Goette (2007) that yields a one-parameter modification of the rule given by (12). In this model, loss averse drivers treat a nominal lease cost of L as if this is really equal to κL for $\kappa > 1$. As in Andersen et al. (2014), our model of loss aversion postulates a time-varying reference point. In this case, it seems natural to assume that the potential earnings that would be realized under the default Rideshare contract determine the reference point for Taxi contracts. Drivers are averse to buying a Taxi contract that ends up reducing their earnings. This produces a kink in the utility of earnings when farebox crosses the Taxi breakeven.

Parametric Lease Aversion

Loss aversion isn’t *necessary* to explain lease aversion, but it does fit the facts.¹⁷ The lease aversion hypothesis is evaluated in combination with a model that describes how drivers predict their earnings. Our parametric forecasting model supposes that driver i ’s forecast of his potential farebox, $y_{0i} = wh_{0i}$, is drawn from a log Normal distribution. Specifically,

¹⁷Chetty and Szeidl (2016) show that consumption commitments can also make moderate stakes gambles unattractive.

conditional on driver characteristics, X_i , forecast y_{0i} is assumed to be distributed according to:

$$\ln y_{0i}|X_i \sim N(X_i'\beta, \tau_0^2), \quad (17)$$

where X_i includes opt-in week and/or earlier earnings and our experimental stratification variables. Using this and (9), the probability driver i participates in Taxi when offered (L_i, t_i) can be written:

$$\begin{aligned} q_0(L_i, t_i; X_i) &= 1 - \Phi \left[\frac{\ln \frac{L_i}{t_i} + \ln \kappa - \sigma(t_i) - X_i'\beta}{\tau_0} \right] \\ &= \Phi \left[\frac{1}{\tau_0} \left(\sigma(t_i) + X_i'\beta - \ln \frac{L_i}{t_i} \right) - \frac{1}{\tau_0} \ln \kappa \right], \end{aligned}$$

where κ is the behavioral lease rate, $\sigma(t_i)$ is the proportional participation threshold reduction due to higher Taxi wages, and (L_i, t_i) describes the Taxi contract offered to this driver.¹⁸ Again, $\sigma(t_i)$ is computed using $\delta^f = 1.8$.

Assuming forecasts are correct on average, β is identified from a regression of log farebox on X_i in the control sample. The parameters of primary interest in this model, τ_0 and κ , can then be estimated by inserting the regressor,

$$\hat{w}_i = \hat{\sigma}(t_i) + X_i'\hat{\beta} - \ln \frac{L_i}{t_i},$$

into a probit model for take-up. Specifically, probit regressions of individual driver participation decisions on \hat{w}_i and a constant identify τ_0 and κ as transformations of the slope and intercept in this conditional probability function

$$P[D_i = 1|L_i, t_i, X_i] = \Phi \left(\frac{1}{\tau_0} \hat{w}_i - \frac{1}{\tau_0} \ln \kappa \right). \quad (18)$$

This model allows forecast earnings variance to exceed the empirical earnings variance. The extra variance in forecast earnings can be motivated by the fact that many drivers will not have known their full opt-in week farebox at the time they decided to lease (most participation decisions were made shortly after receiving the first communication presenting a Taxi offer).

We start with a version of (18) implemented by regressing the log of opt-in week farebox (for control drivers) on a set of covariates, X_i , that includes lags farther back. As can be seen in the first column of Table 7, the resulting estimate of κ is about 1.4, with an

¹⁸For example, a driver in the t_0 fee class at Uber who was offered a zero fee has $\sigma(t_i) = \ln \left[1 + \frac{\delta^f t_0}{2(1 - t_0)} \right]$

estimated forecast standard deviation roughly twice as large than the root mean-squared error (RMSE) of the forecasting regression, (17). The theoretical appendix shows that $\kappa = 1.4$ implies a coefficient of loss aversion around 2, not far from estimates reported in Tversky and Kahneman (1991).

Columns 2-4 of Table 7 report estimates from models incorporating a forecasting equation that predicts farebox during the week Taxi drivers exploited their lease. Columns 2, 3 and 4, respectively, report the results of adding one, two, and then three further farebox lags to the list of predictors in X_i .¹⁹ The resulting estimates of κ , shown along with bootstrapped standard errors computed as described in the empirical appendix, are remarkably stable at around 1.4 in all specifications. Estimates of the standard deviation of the forecast distribution are again quite a bit larger than the RMSE of the corresponding forecasting variance. These estimates suggest that driver uncertainty indeed includes an idiosyncratic component beyond the conditional cross-sectional variance of earnings. At the same time, this extra uncertainty is insufficient to rationalize Taxi undersubscription.²⁰

Nonparametric Lease Aversion

Control drivers' earnings are sampled from the distribution of y_{0i} . The extent of driver lease aversion is therefore identified without need of a parametric model for y_{0i} . To see this, note that incorporating lease aversion in the participation rule given by (9), drivers buy a Taxi lease if

$$\ln y_{0i} > \ln \frac{L}{t} + \ln \kappa - \sigma(t),$$

for any distribution of $\ln y_{0i}$. Writing p_{Lt} for the Taxi participation rate among drivers offered $[L, t]$, this rule implies

$$1 - p_{Lt} = F_0\left(\ln \frac{L}{t} + \ln \kappa - \sigma(t)\right),$$

¹⁹The forecasting models used here include indicators for each of the eight hours \times fee \times week strata. Lag coefficients vary by week offered Taxi. Lagged log earnings are set to zero when lagged earnings are zero; models include missing data dummies for this occurrence. The 2nd lagged farebox for the first Taxi trial includes data from the week in which Wave 2 was driving fee free. The 3rd lagged farebox for the first Taxi trial includes data from the week in which Wave 1 was driving fee free. Note that this model has both a slope (the coefficient on \hat{w}_i) and an intercept. Parameter τ is computed as the reciprocal of the slope. Parameter κ is computed using this estimate and the value of the intercept.

²⁰Table A8 in the Online Appendix shows that there is little heterogeneity in the estimates of κ across subgroups. In particular, we find no evidence that drivers who drive more are less loss averse.

where F_0 is the control drivers' log farebox distribution. Distribution function F_0 can then be inverted to produce a quantile regression that identifies κ :

$$\underbrace{F_0^{-1}(1 - p_{Lt})}_{\text{non-participation quantile}} = \ln \kappa + \ln \frac{L}{t} - \sigma(t) \quad (19)$$

The dependent variable here is the non-participation quantile for the sample of drivers offered $[L, t]$, that is, the farebox value that has p_{Lt} of drivers above and $1 - p_{Lt}$ of drivers below it.

Figure 9 plots the sample analog of $F_0^{-1}(1 - p_{Lt})$ against $\ln \frac{L}{t} - \sigma(t)$ for our 16 Taxi treatment combinations. Without lease aversion (i.e., $\kappa = 1$), the quantiles plotted on the y-axis should be close to the log breakeven minus an adjustment for driver response to higher Taxi wages ($\sigma(t)$), with deviations from this value due solely to sampling variance. The black line in the figure is the forty-five degree line marking these points. As can be seen in the figure, however, non-participation quantiles systematically exceed the adjusted log breakeven. The average gap between predicted and treated quantiles is summarized by the blue regression line, which has slope equal to that generated by a weighted regression of non-participation quantiles on $\ln \frac{L}{t} - \sigma(t)$, with weights given by the number of treated drivers in each hours stratum. Although the estimated slope here is close to one, the empirical quantiles are clearly shifted up, implying that drivers typically set a bar higher than the theoretical breakeven when deciding to buy a Taxi lease.

The intercept generated by the blue line in the figure implies a value of κ equal to about 1.6 (that is, $e^{.45}$). This estimate is similar to those from the parametric model of Taxi take-up, though considerably less precise. Whiskers in the figure denote 95% confidence intervals, computed using bootstrapped standard errors.²¹ Because the non-parametric estimates are less precise than the parametric, parametric estimates are employed in the CV calculations discussed below.

5.2 Accounting for Inattention

As in the Mas and Pallais (2018) analysis of worker response to various sorts of job offers, a simple alternative to the lease aversion story is driver inattention to the details of Earnings Accelerator lease offers. Perhaps some drivers failed to notice or understand our proffered

²¹These are calculated by drawing bootstrap samples of treated and control drivers, stratifying by commission, treatment group, and week, and estimating κ nonparametrically for each bootstrap sample.

Taxi contracts.

As noted above, the Earnings Accelerator promotion was deployed in phases in part to identify drivers attentive to Uber’s extensive promotional messaging. In particular, only drivers who opted to participate in the “no-lose” free week phase were offered Taxi contracts. But some participating drivers may nevertheless have missed or ignored follow-up offers. This in turn provides an alternative explanation for Taxi under-subscription.

We model inattentive behavior by modifying participation equation (18) to include a fraction that ignore promotion. Letting this fraction be denoted by ϕ generates a theoretical participation probability of

$$(1 - \phi)\Phi \left[\frac{1}{\tau_0} \left(\sigma(t_i) + X_i'\beta - \ln \frac{L_i}{t_i} \right) - \frac{1}{\tau_0} \ln \kappa \right]. \quad (20)$$

Inattention can be distinguished from lease aversion because the former is modeled as a fixed proportion of behavioral take-up, while the latter is additive and inside the probit function, with an effect that implicitly depends on covariates.

Columns 5 and 6 of Table 7 present estimates of two versions of this augmented model, one where ϕ is constrained to be equal for all drivers (shown in column 5) and one where ϕ varies with baseline hours (shown in column 6). Results in both columns show little evidence of inattention among participating drivers. Moreover, the estimates of κ generated by models allowing for inattention differ little from those generated by the simpler model. Note also that a nonparametric version of these findings amounts to the claim that the non-participation quantile on the left-hand side of (19) should be shifted additively from the surplus-adjusted breakeven. Figure 9 indeed seems consistent with this hypothesis. Finally, results not reported here explore specifications allowing the probability a worker is attentive to vary as a function of gender and experience on the Uber platform, as well as hours worked. These are consistent with the estimates in Columns 5 and 6 of Table 7.

6 Compensating For Leasing

We use estimates of the ISE and lease aversion coefficient to compute average weekly CV for the sample of 19,316 active Boston drivers described in column 1 of Table 1. This sample drives more and therefore has higher weekly earnings than the sample of eligible drivers, which

is limited to those with average weekly hours between 5 and 25. Conditional on driving, the average weekly farebox in the Boston active sample is \$541 in July 2016; weekly earnings are about \$423. This exceeds the average farebox (and earnings) in Table 1 because here the average is over weeks rather than over drivers and omits weeks with zero earnings. Dropping zeros sidesteps the issue of how or whether to compensate inactive drivers who buy a lease. On one hand, we might assume that future inactive drivers neither drive Rideshare nor lease, in which case their CV is zero; alternately, as in our experiment, inactive drivers who buy a lease might be presumed to be stuck with it, in which case their CV should equal the lease price. The CV calculation for active drivers uses equation 4, with $\delta^f = 1.2$ and $\kappa = 1.4$, which are representative of our findings for the full sample.

Table 8 shows average weekly CV computed for a range of possible Rideshare-Taxi wage gaps and leases. We interpret wage gaps as generated by Rideshare fees, though we can likewise see these gaps as reflecting fare differences under alternative transportation regulations. The CV measures the amount of money required to make a driver as well off under Taxi-style leasing (“Taxi”) as they are under the proportional compensation scheme (“Rideshare”). As can be seen in Panel A of Table 8, for weekly lease rates in the range of the 2010 Boston lease cap of \$700, the average compensation needed to make a driver indifferent between Rideshare and Taxi ranges from \$166 with $L = 600$ and a wage difference of 50%, to \$710 when $L = 800$ and the wage gap is only 15%.²²

With a 25% fee and a lease cost of \$600, perhaps a realistic scenario, average CV is \$437. Almost all active drivers have positive CV in this case (recall that negative CV indicates drivers prefer Taxi; the third entry in each cell indicates the proportion of drivers who prefer Rideshare to leasing). About ten percent of drivers who bought a Taxi contract did not drive in the week covered by their lease. These drivers presumably meant to drive when they bought the lease, but were precluded from doing so, perhaps for reasons related to health or family.²³

With lower lease costs, CV is naturally smaller; in low-lease scenarios, Taxi may well be a better deal. For a lease rate of \$150, for example, a wage gap of 25% makes leasing attractive

²²CV accounts only for the value of the compensation contracts. Because drivers value other characteristics of rideshare, this is an underestimate of the amount drivers would have to be compensated if rideshare opportunities disappeared (Chen et al., Forthcoming, 2020).

²³Lease payments were not refunded to drivers who did not drive.

to many, with average CV equal to -\$13, though 59% still prefer Rideshare in this scenario (median CV is \$35; medians are reported in the second row of each cell in Table 8). With a lease price of only \$100 and a fee of 20%, most drivers (55%) prefer Taxi.

A natural summary measure of CV is the lease price that sets CV equal to 0, that is, the lease rate that leaves drivers indifferent between Rideshare and Taxi. As can be seen in column 8, this averages from \$90 with a 15% wage gap up to \$434 with a 50% wage gap. These maximum lease values are equal to the (average of the) sum of the fees that would be paid to Uber without leasing plus the surplus generated by higher Taxi wages. In the notation of equation (4), this quantity is $t_0wh_0(1 + \frac{\delta t_0}{2(1-t_0)})$.

Behavioral lease values are calibrated to be 40% above nominal (since $\kappa = 1.4$), with the resulting CV calculation summarized in Panel B of Table 8. Assuming Rideshare fees of 25% or less, lease aversion makes CV positive even for a lease cost of only \$150: the Rideshare contract in this case generates average surplus of \$47, though 50% higher Taxi wages make Taxi attractive to most drivers (38% still prefer Rideshare in this case). A lease of \$116 equates Rideshare and Taxi with a 25% fee differential. Even with a lease as low as \$100, however, most lease averse drivers prefer Rideshare to Taxi given a 25% fee gap.

As can be seen in column 5 of Panel B of Table 8, with a \$400 lease and a 25% wage difference, median CV is \$445, more than the nominal lease. The excess of CV over the nominal lease can be interpreted as an interest payment to drivers in return for lending the local Taxi and Limousine Commission (or other lease-granting authority) the value of the lease until compensation is paid (presumably 1-2 weeks after lease purchase, that is, the week after leased driving is completed). Interest of \$45 on a \$400 loan for a week or two sounds high, but not out of line with the \$15 fee typically paid for a \$100 payday loan.²⁴

The comparisons in Table 8 implicitly make driving Taxi a condition for receipt of compensation. An alternative compensation scenario allows former Rideshare drivers to quit driving when Rideshare work disappears, receiving UI instead (this is fanciful since rideshare drivers who stop driving don't currently qualify for UI). As noted at the end of Section 2.2, UI cuts the monetary cost of compensation by allowing former Rideshare drivers to take additional leisure. The dollar compensation required to make idle Rideshare drivers as well

²⁴The cost of payday loans is described in <http://libertystreeteconomics.newyorkfed.org/2015/10/reframing-the-debate-about-payday-lending.html>.

off as they were when driving for Rideshare is reported in Appendix Table A9, along with the proportion expected to take this option.

Appendix Table A9 shows that the UI option reduces the cash compensation required to make former rideshare drivers indifferent to the disappearance of the rideshare compensation scheme. Importantly, however, the UI compensation option also slashes consumer (rider) welfare by reducing the supply of drivers. With a \$200 lease and a 25% wage difference for example, 48% of non-lease-averse drivers take advantage of the opportunity to receive compensation without driving (the proportion sitting out appears in the second line of each cell). This reduces the number of hours supplied to the market by 17% (these figures appear in the third row of each cell).

In the UI version of this scenario with lease averse drivers, UI cuts service by almost a third. Because CV-compensated drivers are necessarily just as well off either way, the contrast that requires Taxi driving as a condition for compensation comes closer to a welfare comparison focused on drivers, while leaving rider welfare improved or unchanged (in fact, the driving requirement weakly increases trip supply). The non-UI scenario is also fiscally attractive: in principle, a benevolent Taxi and Limousine Commission can implement the scheme described in Table 8 using the revenue from leasing, with some money left over. Historically, however, the revenue from medallion sales has not been redistributed to drivers. It's also worth noting that a long-term, unanticipated removal of rideshare work opportunities may have income effects, meaning the relevant elasticity for welfare calculations is smaller. A smaller labor supply elasticity makes Taxi less attractive, increasing the compensation required when rideshare work disappears.²⁵

7 Summary and Directions for Further Work

Compensation schemes are integral to work arrangements of many kinds. A key parameter governing the economic consequences of alternative work arrangements is the response of work effort to pay. A high labor supply elasticity makes lease-based contracts more attractive since

²⁵CV is also larger if labor supply is less elastic to the ride-hailing industry as a whole than to individual platform operators (Caldwell and Oehlsen, 2018). The distribution of total rideshare hours worked may differ from the distribution of Uber hours. But in our data, Uber drivers appear to earn over 90% of their ride-hailing income from Uber (Koustas, 2017). This suggests our CV calculation for a sample of Uber drivers is not too wide of the mark.

workers that are more responsive to pay gain more from higher wages. On the other hand, the kind of lease aversion documented here, a context-specific form of loss aversion, increases the compensation needed to induce leasing.

Our randomized Taxi experiment identifies an ISE on the order of 1.2 in the full sample and 1.8 for drivers who opted in to leasing. These values are not large enough to overcome many drivers’ lease aversion. Consequently, the compensation required to make drivers indifferent to the loss of rideshare-type earnings opportunities far exceeds the already mostly-positive CV computed using nominal lease rates. Perhaps surprisingly, the response to wage changes and lease offers varies little across strata defined by hours worked and commission rates (which vary with experience). This suggests the estimates reported here may be reasonably representative of Uber driver behavior.

Our estimates of the value of the proportional compensation model come from Uber drivers—who have selected into this type of work. Evidence for lease aversion in other settings comes from the New York City TLC, which recently began piloting “Fare Share Leasing” allowing drivers to lease a medallion by paying “a percentage of a driver’s farebox revenue”, much like the Uber fee.²⁶ This pilot is in response to the results of a survey the TLC conducted coincident with the biennial review of fares and lease caps; one of the driver issues identified in this survey was the “stress associated with starting shifts ‘in the red’ having paid a set lease price at the beginning of shifts”. Interest in alternatives to leasing for taxi drivers may be boosted by technological change, such as electronic payments and app-based dispatching technology, that limits the scope for drivers to under-report earnings.

Our economic analysis focuses on drivers. In principle, however, the experimental Taxi scheme evaluated here creates enough additional surplus to allow drivers and platform owners to negotiate a lease-based contract that is Pareto superior to commission-based compensation schemes like the Uber fee. As is the case with any system that taxes output, the social cost of the Rideshare contract arises from the wedge proportional fees insert between wages and effort. Medallion leasing effectively “sells the firm to the worker,” a classic solution to the problem of efficient contracting (see, e.g., Lazear 1995, 2018). Lease revenue is therefore adequate to compensate drivers who might prefer relatively efficient Taxi-style contracts. But this compensation possibility presumes drivers will indeed accept nominal CV in return

²⁶See https://www1.nyc.gov/assets/tlc/downloads/pdf/taxicab_leasing_resolution.pdf for details.

for leasing.

Our results hint at why the rapidly evolving ride-hailing market seems to have only briefly flirted with virtual leasing of the sort explored in our experiment. In 2016, Boston rideshare upstart Fasten offered its drivers an \$80 lease in return for “weekly unlimited driving,” that is, driving with no fee. Fasten’s other compensation scheme took a fee equal to a dollar a trip; this probably amounts to an average fee around 10%. As can be seen in Panel B of Table 8, with a 15% fee, any lease under \$90 is attractive. Fasten’s \$80 lease therefore seems likely to have been in the ballpark for many drivers. But this conclusion is overturned by lease aversion, which reduces the maximum lease rate that drivers will pay to avoid a 15% fee to \$64. It’s unsurprising, therefore, that Fasten appears to have had few takers for weekly unlimited driving.²⁷

Finally, it’s interesting to compare our results with the Mas and Pallais (2018) estimates of workers’ willingness-to-pay for job amenities. Their findings suggest workers place little value on hours flexibility for its own sake, though they prefer to avoid jobs that grant employers discretion when setting work schedules. These findings seem consistent with the notion that it’s the need to pay lease costs up front rather than hours constraints that make leasing distasteful. Looking down the road, a natural direction for future research is an exploration of how the economic consequences of contractual differences interact with market structure, such as the presence of competing ride-hailing services. More competition between service providers presumably means a more elastic labor supply response to individual platform operators. This in turn should make Taxi contracts like ours more attractive.

²⁷Fasten ended operations in the United States in early 2018.

Table 1: Boston Uber Drivers

	All Boston Drivers (1)	Eligible Drivers (2)	Experimental Drivers (3)	Strata-Adjusted Difference (4)
Female	0.14 (0.35)	0.14 (0.34)	0.14 (0.35)	0.00 (0.01)
Age	40.90 (12.13)	41.58 (12.20)	41.80 (12.29)	0.15 (0.36)
Hours Week Starting 08/08	19.58 (15.91)	15.14 (9.17)	16.24 (9.23)	-0.05 (0.26)
Average Hours/Week Month Before Selection	14.42 (14.39)	13.13 (5.69)	14.51 (5.81)	0.06 (0.08)
Average Hourly Earnings Month Before Selection	15.39 (8.64)	17.59 (6.19)	17.40 (6.05)	-0.10 (0.17)
Average Weekly Farebox Month Before Selection	372.06 (447.51)	310.91 (192.04)	342.82 (198.12)	-0.80 (3.93)
Months Since Sign-up	13.89 (9.43)	14.26 (9.25)	11.14 (8.67)	-0.08 (0.15)
Vehicle Solutions	0.08 (0.26)	0.08 (0.27)	0.08 (0.28)	0.01 (0.01)
Car Model Year 2003 or Older	0.03 (0.17)	0.03 (0.17)	0.12 (0.33)	--- ---
Car Model Year 2011 or Newer	0.64 (0.48)	0.64 (0.48)	0.56 (0.50)	-0.01 (0.01)
Commission	22.34 (2.50)	22.24 (2.49)	23.21 (2.40)	0.00 (0.01)
Number of Observations	19316	8685	1600	8685

Note: Columns 1-2 compare Boston drivers to the subset of drivers eligible for the experiment. Eligible drivers are those with valid vehicle year information who made at least 4 trips during the past 30 days and drove an average of between 5 and 25 hours/week in July 2016. Column 3 shows means for treated drivers. Treatment was randomly assigned within strata defined by hours (high/low), commission (20/25% commission) and car age (older/newer than 2003). Column 4 shows the strata-adjusted difference between the treated sample and the eligible pool. Average hourly earnings include surge but are net of fee. Vehicle solutions drivers lease a car through an Uber-sponsored leasing program.

Table 2: Earnings Accelerator Opt-In Week Design and Take-up

Group			Offers		Opt-Ins	
Hours	Car	Fee	Number	Rate	Number	Rate
(1)	(2)	(3)	(4)	(5)	(6)	(7)
<u>Wave 1</u>						
High	New	20%	102	6%	75	74%
	New	25%	202	17%	155	77%
	Old	--	96	100%	61	64%
		--	400	13%	291	73%
Low	New	20%	100	4%	68	68%
	New	25%	200	8%	148	74%
	Old		100	54%	64	64%
		--	400	7%	280	70%
Total			800		571	
<u>Wave 2</u>						
High	New	20%	150	8%	84	56%
	New	25%	250	21%	154	62%
		--	400	13%	238	60%
Low	New	20%	250	9%	133	53%
	New	25%	150	6%	89	59%
		--	400	7%	222	56%
			800		460	

Note: This table describes the initial pair of “opt-in weeks”, designated wave 1 and 2. This phase of the experiment randomly assigned offers within the groups defined by columns 1-3: hours bandwidth, car age, and commission. Column 4 shows the number of offers that were made to this group. Column 5 shows the percent of eligible Boston drivers in that hour and commission group that were offered fee-free driving. Column 6 reports the number of drivers who accepted the offer and column 7 reports the percentage of drivers who accepted the offer. The last row in each hours group shows the total for that week and hours group; the last row in each wave shows the overall total for that week.

Table 3: Who Opts In?

	Pooled		High Hours		Low Hours	
	Opt-Out		Opt-Out		Opt-Out	
	Mean	Difference	Mean	Difference	Mean	Difference
	(1)	(2)	(3)	(4)	(5)	(6)
Female	0.13	0.03	0.12	0.01	0.14	0.05*
	[0.33]	(0.02)	[0.32]	(0.02)	[0.35]	(0.03)
Age	42.75	-1.46**	44.81	-3.06***	40.89	-0.08
	[12.61]	(0.65)	[12.68]	(0.95)	[12.27]	(0.89)
Commission	23.11	0.16	22.97	0.33*	23.24	0.00
	[2.43]	(0.13)	[2.46]	(0.18)	[2.39]	(0.17)
Vehicle Solutions	0.06	0.03**	0.07	0.04**	0.05	0.02
	[0.24]	(0.01)	[0.26]	(0.02)	[0.23]	(0.02)
Vehicle Year	2010.40	-1.76	2010.56	-3.50	2010.26	0.06
	[4.45]	(1.96)	[4.39]	(3.82)	[4.51]	(0.33)
Months Since Signup	11.60	-0.71	12.53	-1.61**	10.75	0.09
	[9.03]	(0.46)	[9.19]	(0.67)	[8.81]	(0.63)
Hours Worked Week Starting 08/22	11.28	4.01***	16.07	2.56**	6.93	4.86***
	[13.35]	(0.69)	[14.48]	(1.06)	[10.51]	(0.79)
Farebox Week Starting 08/22	251.50	99.93***	358.77	71.83***	153.95	114.05***
	[306.38]	(16.07)	[340.60]	(25.15)	[232.39]	(17.91)
Average Hours/Week the Month Before Selection	14.16	0.53*	19.67	-0.18	9.16	0.49**
	[6.01]	(0.31)	[3.01]	(0.22)	[2.84]	(0.21)
Average Hourly Earnings the Month Before Selection	16.19	1.88***	17.46	2.30***	15.03	1.26***
	[5.56]	(0.30)	[4.80]	(0.38)	[5.95]	(0.45)
Average Weekly Farebox the Month Before Selection	310.06	50.85***	447.65	57.13***	184.93	24.36***
	[180.52]	(9.90)	[145.18]	(11.48)	[100.90]	(7.54)
Number of Observations	569	1600	271	800	298	800

Note: This table compares the characteristics of drivers who opted-in to fee-free driving with those of drivers who were offered fee-free driving but did not participate. Standard deviations appear in brackets. Columns 2, 4, and 6 report the strata-adjusted difference between drivers who opted-in and drivers who did not opt-in. Standard errors are in parentheses. Average hourly earnings include surge but are net of fee. Vehicle solutions drivers lease a car through an Uber-sponsored leasing program. Levels of significance: *10%, ** 5%, and *** 1%.

Table 4: Earnings Accelerator Taxi Parameters and Take-up

Group			Treatment			Offers and Opt-Ins	
		No. in			Break-	Offer	Opt-In
Hours	Fee	Group	Lease	New Fee	even	Rate	Rate
(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
<u>Week 1</u>							
High	0.20	180	\$110	0	\$550	0.4	0.42
			\$165	-0.125	\$508	0.2	0.53
High	0.25	349	\$110	0	\$440	0.4	0.28
			\$165	-0.125	\$440	0.2	0.33
Low	0.2	177	\$45	0	\$225	0.4	0.58
			\$75	-0.125	\$231	0.2	0.51
Low	0.25	325	\$45	0	\$180	0.4	0.48
			\$75	-0.125	\$200	0.2	0.34
<u>Week 2</u>							
High	0.20	180	\$60	0	\$300	0.3	0.50
			\$25	0.10	\$250	0.3	0.46
High	0.25	349	\$55	0	\$220	0.3	0.41
			\$35	0.125	\$280	0.3	0.54
Low	0.2	177	\$40	0	\$200	0.3	0.43
			\$15	0.10	\$150	0.3	0.58
Low	0.25	324	\$35	0	\$140	0.3	0.43
			\$15	0.125	\$120	0.3	0.58

Note: This table describes the “Taxi” phase of the experiment. Columns 1 and 2 define the strata (hours group and commission) used in phase two. During each of two Taxi weeks, drivers within each stratum were randomly assigned to one of two treatments (60%) or to the control group (40%). Columns 4 and 5 describe these treatments. Column 6 reports the break-even associated with each treatment. Opt-in rates in column 8 are reported as a proportion of drivers offered. Lease prices were chosen so as to be attractive to roughly sixty percent of the drivers in each stratum.

Table 5: Estimated ISEs

	Opt-In Week			Taxi		
	Pooled	High Hours	Low Hours	Pooled	High Hours	Low Hours
	(1)	(2)	(3)	(4)	(5)	(6)
A. 2SLS Estimates						
First Stage	0.20*** (0.01)	0.19*** (0.01)	0.21*** (0.02)	0.13*** (0.02)	0.11*** (0.02)	0.15*** (0.02)
2SLS	1.16*** (0.12)	1.12*** (0.16)	1.21*** (0.19)	1.81*** (0.37)	2.18*** (0.66)	1.49*** (0.44)
Over-identified Model	1.17*** (0.12)	1.12*** (0.16)	1.23*** (0.19)	1.48*** (0.27)	1.46*** (0.40)	1.54*** (0.39)
B. OLS Estimates						
OLS	0.21*** (0.06)	0.13* (0.08)	0.29*** (0.09)	0.03 (0.08)	-0.05 (0.09)	0.13 (0.14)
Drivers	1176	649	527	822	445	377
Observations	2214	1242	972	1422	775	647

Note: This table reports 2SLS estimates of the intertemporal substitution elasticity (ISE). The endogenous variable is log wages, instrumented with dummies indicating treatment offers. The over-identified model in columns 1-3 uses separate treatment indicators for each week, fee class, and hours group. The over-identified model in columns 4-6 uses separate treatment indicators for each taxi offer. These offers vary by week, treatment group, fee class, and hours group. Models control for the strata used for random assignment, time dummies, gender, whether a driver uses Uber's vehicle solutions program, the number of months since sign-up, whether the car is older than 2003, and one lag of log earnings. Standard errors are clustered by driver. A total of 1600 drivers were offered fee-free driving in opt-in week; 1031 accepted the offer and were eligible for Taxi leasing. Sample sizes in columns 1 and 4 are lower because the data used to construct this table omit zeros. Levels of significance: *10%, ** 5%, and *** 1%.

Table 6: Gains and Losses from Taxi

	Opt-In Week Earnings		Live Week Earnings	
	Expected		Observed	Expected
	Did Not		Participated	Did Not
	Participated	Participate		Opt-In
	(1)	(2)	(3)	(4)
A. All				
Mean Benefit	\$92	\$66	\$85	\$64
Percent with Positive Benefit	78%	56%	85%	54%
Observations	560	679	560	679
B. Conditional on Driving During Live Week				
Mean Benefit	\$103	\$106	\$97	\$115
Percent with Positive Benefit	83%	78%	92%	87%
Observations	515	423	515	423

Note: This table reports the median gains and losses from the Taxi treatment among treated drivers who did and did not buy a taxi contract. Columns 1 and 2 use data from Taxi opt-in week. Columns 3 and 4 use the same data, but adjust driver hours using the experimental wage offer and an ISE of 1.2. Panel A includes data for all treated drivers. Panel B includes data for drivers who drove during live week. The first row in each panel presents the mean gain for all workers in the sample. The second row presents the percent of workers with positive gains. The third row presents the median gain for drivers with positive gains. Levels of significance: *10%, **5%, and *** 1%.

Table 7: Modeling Taxi Take-Up

	Parametric				Inattention	
	(1)	(2)	(3)	(4)	(5)	(6)
Slope	0.69*** (0.10)	0.73*** (0.09)	0.81*** (0.09)	0.79*** (0.09)	0.69*** (0.10)	0.68*** (0.10)
Intercept	-0.24*** (0.07)	-0.25*** (0.07)	-0.28*** (0.07)	-0.27*** (0.07)	-0.24*** (0.07)	-0.17* (0.09)
Implied Kappa	1.41*** (0.11)	1.40*** (0.10)	1.41*** (0.10)	1.41*** (0.10)	1.41*** (0.11)	1.27*** (0.15)
Implied Tau	1.46*** (0.22)	1.36*** (0.18)	1.24*** (0.15)	1.26*** (0.16)	1.46*** (0.22)	1.47*** (0.23)
Forecasting regression RMSE	0.71	0.82	0.80	0.79	0.71	0.71
Attentive					1.00*** (0.00)	
Attentive * Low Hours						0.91*** (0.06)
Attentive * High Hours						1.00*** (0.01)
Number of Drivers	954	938	938	938	954	954
Earnings Distribution	Predicted Opt-In Week	Predicted Live Week	Predicted Live Week	Predicted Live Week	Predicted Opt-In Week	Predicted Opt-In Week
Number of Earnings Lags	1	1	2	3	1	1
Covariates	Yes	Yes	Yes	Yes	Yes	Yes


Notes: Parametric models are fit to micro data on take-up using equation (18) in the text. Standard errors are bootstrapped as described in the appendix. Each column uses data from the control drivers' earnings distribution. Note that, because κ is unobserved, equation (18) has both a slope (on \hat{w}_i) and an intercept. We obtain our estimates of τ by taking the reciprocal of the slope. We obtain our estimates of κ using this estimate and the value of the intercept. To ensure the same sample is used across columns we replace lags of 0 earnings with 0, and include corresponding dummy variables. Appendix Table A8 reports similar results, broken out by subgroup. Levels of significance: *10%, ** 5%, and *** 1%.

Table 8: Compensating Variation for Loss of Rideshare

Wage Gap	Weekly Lease Rates							Max Lease
	\$50 (1)	\$100 (2)	\$150 (3)	\$200 (4)	\$400 (5)	\$600 (6)	\$800 (7)	
A. Nominal Lease								
15%	-\$40	\$10	\$60	\$110	\$310	\$510	\$710	\$90
	-\$13	\$37	\$87	\$137	\$337	\$537	\$737	
	42%	66%	80%	89%	99%	100%	100%	
20%	-\$75	-\$25	\$25	\$75	\$275	\$475	\$675	\$125
	-\$38	\$12	\$62	\$112	\$312	\$512	\$712	
	33%	55%	69%	79%	97%	100%	100%	
25%	-\$113	-\$63	-\$13	\$37	\$237	\$437	\$637	\$163
	-\$65	-\$15	\$35	\$85	\$285	\$485	\$685	
	26%	46%	59%	70%	91%	98%	100%	
50%	-\$384	-\$334	-\$284	-\$234	-\$34	\$166	\$366	\$434
	-\$256	-\$206	-\$156	-\$106	\$94	\$294	\$494	
	10%	20%	29%	37%	59%	74%	83%	
B. Behavioral Lease								
15%	-\$20	\$50	\$120	\$190	\$470	\$750	\$1,030	\$64
	\$7	\$77	\$147	\$217	\$497	\$777	\$1,057	
	54%	78%	90%	96%	100%	100%	100%	
20%	-\$55	\$15	\$85	\$155	\$435	\$715	\$995	\$89
	-\$18	\$52	\$122	\$192	\$472	\$752	\$1,032	
	43%	66%	80%	89%	100%	100%	100%	
25%	-\$93	-\$23	\$47	\$117	\$397	\$677	\$957	\$116
	-\$45	\$25	\$95	\$165	\$445	\$725	\$1,005	
	35%	57%	71%	81%	98%	100%	100%	
50%	-\$364	-\$294	-\$224	-\$154	\$126	\$406	\$686	\$310
	-\$236	-\$166	-\$96	-\$26	\$254	\$534	\$814	
	14%	27%	38%	47%	71%	85%	92%	


Notes: Panel A shows compensating variation (CV, paid to Rideshare drivers to induce them to work under Taxi), computed for the nominal lease rates listed in columns 1-7. Column 8 reports the mean lease that makes a driver indifferent between Taxi and Rideshare. Panel B evaluates CV using behavioral lease rates computed from Taxi take-up. The behavioral lease is fifty percent greater than the nominal lease. The ISE is set at 1.2. The first row of each cell shows average CV. The second row shows median CV. The third row reports the proportion of drivers with positive CV. CV is evaluated using weekly earnings and hours data for all Boston Uber drivers in the month of July who completed at least 4 trips. Weeks with zero trips are omitted. The mean farebox conditional on driving is 541. The mean payout conditional on driving is 423.

Figure 1: Rideshare and Taxi Budget Lines




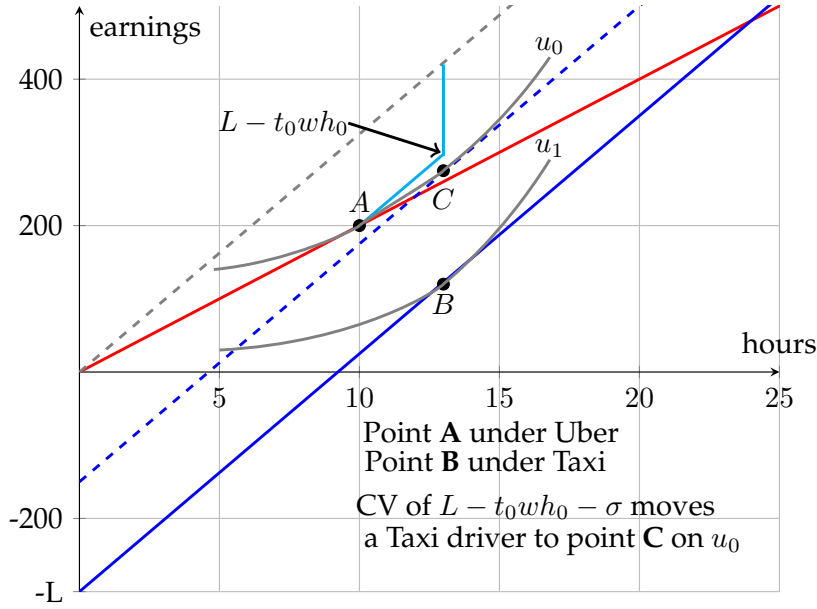
Note: This figure contrasts the Rideshare and Taxi budget sets. The red line shows the Rideshare budget for a driver who collects \$20 in fares each hour and pays a 25% fee. The blue line shows the corresponding Taxi budget set given a lease price of \$100. The two lines cross at the “breakeven” point. The two black lines depict indifference curves for a driver who prefers the proportional fee (u_0) and for a driver who prefers to lease (u_1).

Figure 2: Driven and Elastic Drivers Take Taxi



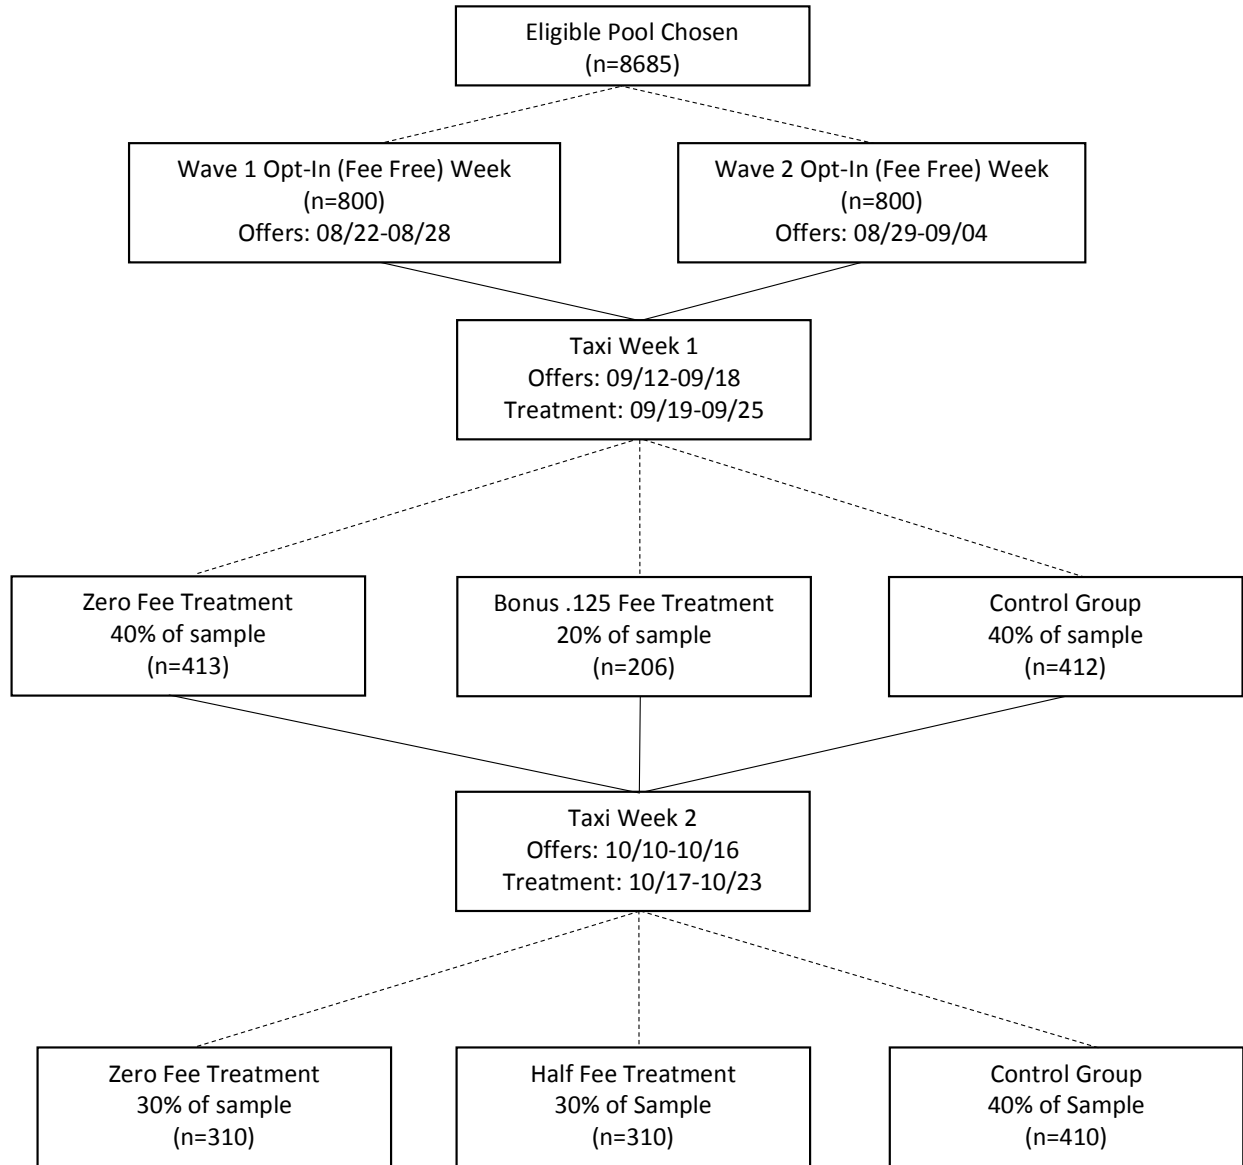
Note: This figure sketches choices made by two types of drivers, those whose Rideshare hours are such that they would earn more by leasing, and those who earn less ex ante but are elastic enough to find leasing attractive anyway. Indifference curves u_0 and u_1 belong to a driver in the second group. While her Rideshare hours h_0 are not sufficient to pay off the lease, she increases her hours worked to h_1 when offered Taxi.

Figure 3: Rideshare vs Taxi Compensating Variation



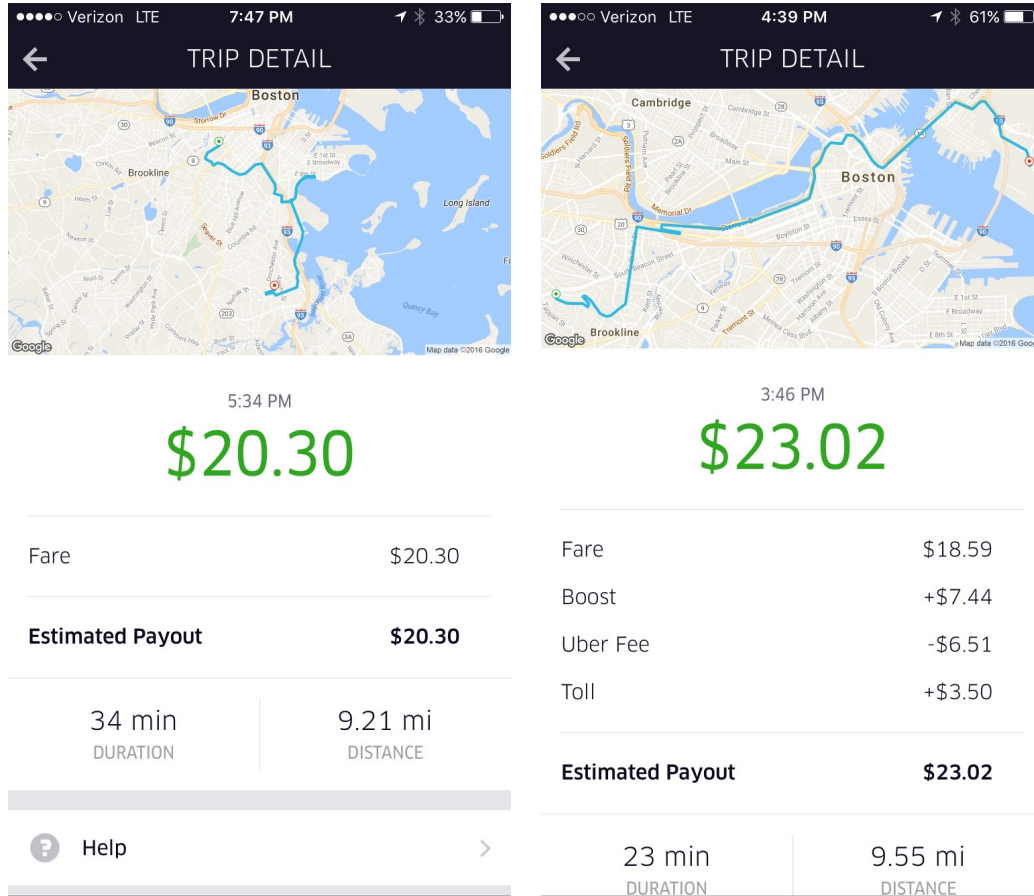
Note: This figure on the left shows how to compute compensating variation for a driver who moves from Rideshare to Taxi. Indifference curve u_0 is tangent to the Rideshare budget set at point A , where this Rideshare driver works 10 hours/week and earns \$200/week. A Taxi contract without compensation moves this driver to point B . Point C indicates the point on u_0 with the same slope as the Taxi budget set. CV for Taxi is given by the distance from B to C , and falls below lease costs net of ex ante fees. The dashed grey line is parallel to the blue line, shifted up by the lease amount, L . The arrow shows the utility level associated with working h_1 if the driver is compensated by the naive amount: the full cost of the lease, less the reduction in fees, calculated using Rideshare hours. The fact that this is above point C —the point along the original indifference curve associated with h_1 —illustrates that this naive compensation level is too high. In order to compensate the driver and move her to point C , the driver must be given $L - t_0wh_0 - \sigma$, where σ is an adjustment term. We describe what this adjustment depends on in section 2.2. The figure on the right shows how this calculation is modified if the driver has access to “UI” of the sort described in section 2.2.

Figure 4: Earnings Accelerator Experimental Design



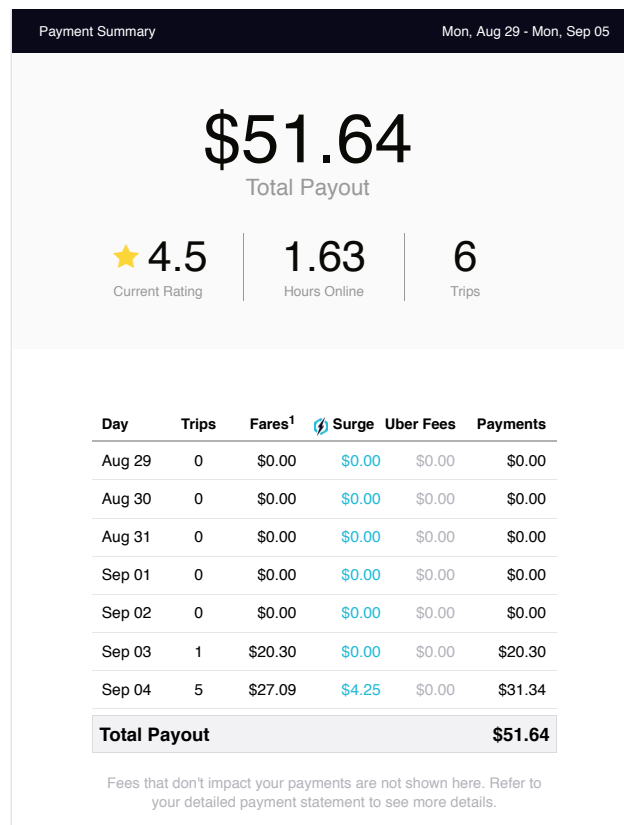
Note: This figure sketches the sequencing and sample sizes associated with each phase of the Earnings Accelerator experiment. Dashed lines denote groupings determined by random assignment.

Figure 5: Earnings Accelerator Trip Receipts




Note: This figure shows pictures of two of a driver's trip receipts. The picture on the left is for a trip taken while the Earnings Accelerator was active; this driver paid no fee. The picture on the right is for a trip taken when the Earnings Accelerator was inactive; this driver paid an Uber fee. The picture on the right shows how additional Uber promotions, such as Surge or Boost (a similar, geocode- and time-specific) are reflected on trip receipts.

Figure 6: Weekly Pay Statements

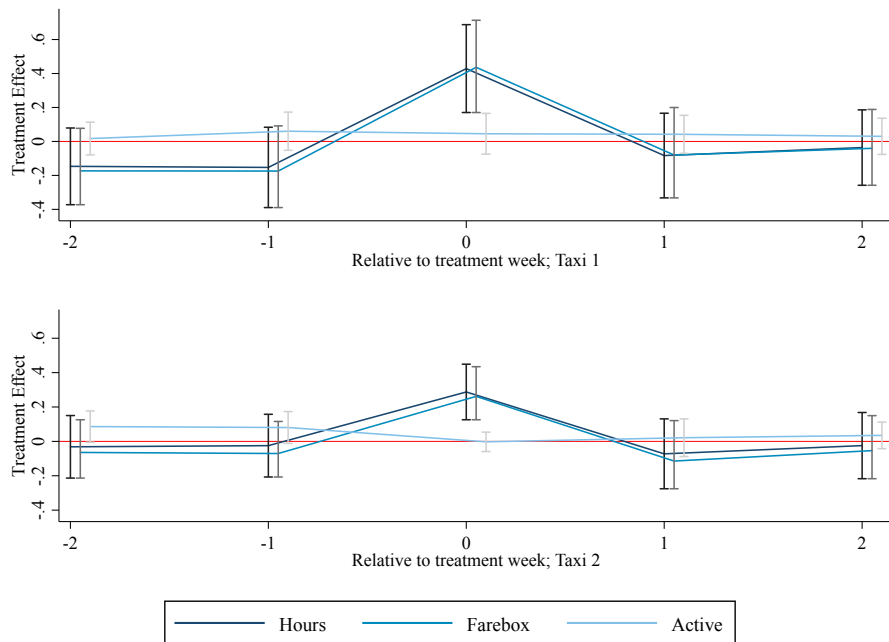


Note: This figure shows a sample weekly pay statement for a driver on the Earnings Accelerator. This pay statement breaks down the drivers' pay into the amount of fares before Surge, the amount earned on Surge, and the amount of fees paid. The column at the right lists the net pay for each day.

Figure 7: Participation Effects on Labor Supply
A. Opt-in Week




B. Taxi




Note: These figures report treatment effects on hours, earnings and an indicator of any Uber activity for drivers who opted in to the Earnings Accelerator. Panel A reports estimates for drivers who accepted the opportunity to drive fee-free. Panel B reports estimates for drivers who bought a Taxi lease. Effects are computed by instrumenting experimental participation with experimental offers as described in the text. Models control for the strata used for random assignment.

Figure 8: Taxi Under-Subscription



Notes: For each of 16 strata defined by pre-experimental hours driven, treatment week, and Taxi treatment offered, this figure plots empirical Taxi participation (lease purchase) rates against the theoretical rate predicted by the treated groups' earning distributions during opt-in week. Diamonds are used for the high hours group and circles are used for the low hours group. Red points are used for the 20% commission group and blue points are used for the 25% commission group. By construction, drivers in the 20% commission group are more experienced than those in the 25% commission group. The ISE is set at 1.8. The dashed line indicates the locus of equality for theoretical and empirical take-up. Rates are calculated on the sample of drivers who drove during opt-in week.

Figure 9: Comparing Empirical and Theoretical Participation Quantiles



Notes: For each of 16 strata defined by pre-experimental hours driven, treatment week, and Taxi treatment offered, this figure plots the quantile of opt-in week earnings for the control group against the log of theoretical opt-in earnings, defined as breakeven minus a labor supply adjustment. Quantiles are calculated on the sample of drivers who drove during opt-in week. Quantiles are evaluated using empirical participation rates. Whiskers indicate 95% confidence intervals for each quantile. A weighted regression line fit to the plotted points appears in blue. A 45 degree line is plotted in black.

References

- Andersen, Steffen, Alec Brandon, Uri Gneezy, and John A List**, “Toward An Understanding of Reference-Dependent Labor Supply: Theory and Evidence From A Field Experiment,” Working Paper 20695, National Bureau of Economic Research November 2014.
- Angrist, Joshua, Sydnee Caldwell, and Jonathan Hall**, “Uber vs. Taxi: A Driver’s Eye View,” *Unpublished Data*, 2020.
- Ashenfelter, Orley**, “Determining Participation in Income-Tested Social Programs,” *Journal of the American Statistical Association*, September 1983, 78 (383), 517–525.
- Browning, Martin, Angus Deaton, and Margaret Irish**, “A Profitable Approach to Labor Supply and Commodity Demands over the Life-Cycle,” *Econometrica*, May 1985, 53 (3), 503–543.
- Caldwell, Sydnee and Emily Oehlsen**, “Monopsony and the Gender Wage Gap: Experimental Evidence from the Gig Economy,” Working Paper, Massachusetts Institute of Technology 2018.
- Castillo, Juan Camilo**, “Who Benefits from Surge Pricing?,” 2019.
- Chen, M Keith, Judith A Chevalier, Peter E Rossi, and Emily Oehlsen**, “The Value of Flexible Work: Evidence from Uber Drivers,” *Journal. of Political Economy*, Forthcoming.
- , —, —, —, and **Lindsey Currier**, “Suppliers and Demanders of Flexibility: The Demographics of Gig Work,” 2020.
- Chetty, Raj**, “A New Method of Estimating Risk Aversion,” *The American Economic Review*, December 2006, 96 (5), 1821–1834.
- and **Adam Szeidl**, “Consumption Commitments and Habit Formation,” *Econometrica*, March 2016, 84 (2), 855–890.

- Einav, Liran, Amy Finkelstein, Stephen P Ryan, Paul Schrimpf, and Mark R Cullen**, “Selection on Moral Hazard in Health Insurance,” *The American Economic Review*, February 2013, *103* (1), 178–219.
- Farber, Henry S.**, “Is Tomorrow Another Day? The Labor Supply of New York City Cabdrivers,” *Journal of Political Economy*, 2005, *113* (1), 46–82.
- , “Why you Can’t Find a Taxi in the Rain and Other Labor Supply Lessons from Cab Drivers*,” *The Quarterly Journal of Economics*, 2015, *130* (4), 1975–2026.
- Fehr, Ernst and Lorenz Goette**, “Do Workers Work More if Wages are High? Evidence from a Randomized Field Experiment,” *The American Economic Review*, March 2007, *97* (1), 298–317.
- Gronau, Reuben**, “Leisure, Home Production, and Work— The Theory of the Allocation of Time Revisited,” *Journal of Political Economy*, December 1977, *85* (6), 1099–1123.
- Hall, Jonathan V and Alan B Krueger**, “An Analysis of The Labor Market For Uber’s Driver-Partners in The United States,” *Industrial and Labor Relations Review*, forthcoming 2017.
- Koustas, Dmitri**, “Labor Supply and Consumption Insurance: Evidence from Rideshare Drivers,” Working Paper, UC Berkeley 2017.
- Lazear, Edward P.**, *Personnel Economics*, Cambridge, MA: MIT Press, 1995.
- , “Compensation and Incentives in the Workplace,” *Journal of Economic Perspectives*, 2018, *32* (3), 195–214.
- Lazear, Edward P. and Paul Oyer**, “Personnel Economics,” in Robert Gibbons and John Roberts, eds., *The Handbook of Organizational Economics*, Princeton University Press, 2012.
- MaCurdy, Thomas E.**, “An Empirical Model of Labor Supply in a Life-Cycle Setting,” *Journal of Political Economy*, December 1981, *89* (6), 1059–1085.

- Mas, Alexandre and Amanda Pallais**, “Valuing Alternative Work Arrangements,” *American Economic Review*, 2018.
- Otsuka, Keijiro and Naoki Murakami**, “Resource Allocation and Efficiency of Sharecropping Under Uncertainty,” *Asian Economic Journal*, March 1987, 1 (1).
- Rabin, Matthew**, “Risk Aversion and Expected-Utility Theory: A Calibration Theorem,” *Econometrica*, September 2000, 68 (5), 1281–1292.
- Sydnor, Justin**, “(Over)insuring Modest Risks,” *American Economic Journal: Applied Economics*, October 2010, 2 (4), 177–199.
- Tversky, Amos and Daniel Kahneman**, “Loss Aversion in Riskless Choice: A Reference-Dependent Model,” *The Quarterly Journal of Economics*, November 1991, 106 (4), 1039–1061.